EPIC
Institute of Technology
Powered by epam

Operating System Internals

Operating Systems

Alexander Kernozhitsky

EPIC
Institute of Technology
Powered by epam

Contents

EPIC
Institute of Technology
Powered by <epam>

01

CPU and Scheduling

EPIC
Institute of Technology
Powered by epam

CPU

- Execute commands one by one
- React to interrupts (≈ signals)
  - Hardware events
  - Software exceptions (division by zero, page fault, etc.)
  - Double fault, triple fault
- Sleeps until an interrupt is received
- Multi-core, each core executes commands independently

EPIC
Institute of Technology
Powered by epam

Processes

- Up to ~200 CPU cores
- Thousands of processes
- Most processes are almost always idle and only react to some events
- How to schedule processes between cores?
  - Cooperative multitasking (very old Windows)
  - Preemptive multitasking (most modern OSes)
  - No multitasking at all (MS-DOS)

EPIC
Institute of Technology
Powered by epam

Cooperative multitasking

- Each process decides when to give control back to OS
- What if process hangs?

EPIC Institute of Technology
Powered by epam

Preemptive multitasking

- Each process is preempted by timer
- Timer generates interrupt and gives control back to OS

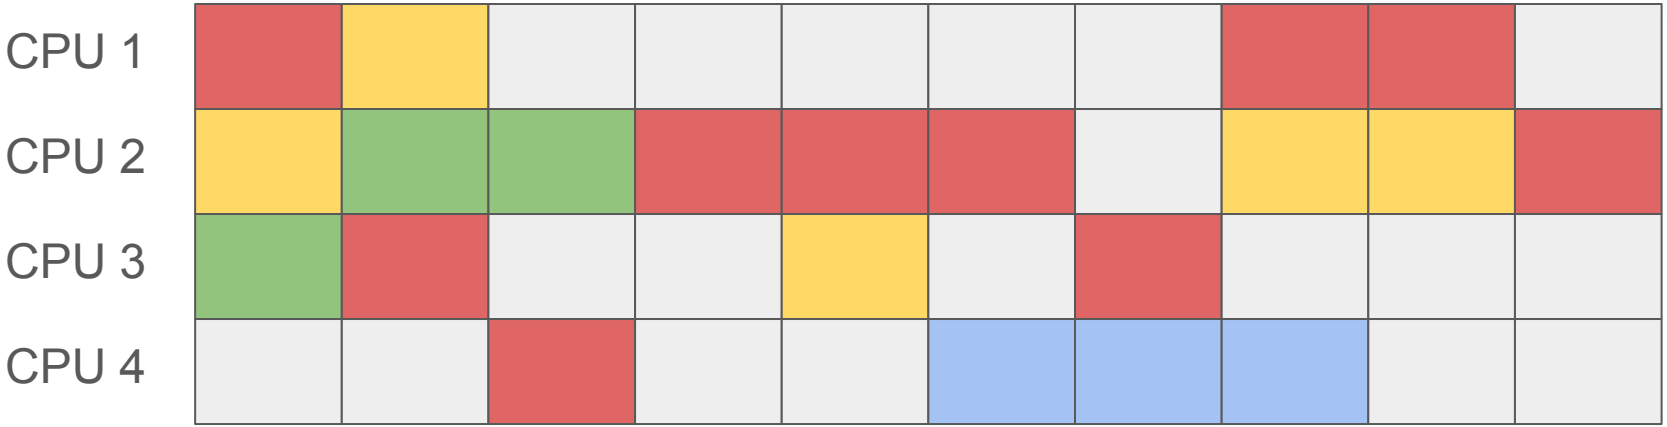

EPIC
Institute of Technology
Powered by epam

Protection

- Protection rings: 0, 1, 2, 3
- Privileged instructions
- Kernel memory

EPIC
Institute of Technology
Powered by <epam>

02

Virtual Memory

EPIC
Institute of Technology
Powered by epam

Memory

- A sequential array of data
- Some parts are reserved for communication with hardware
- How to split between processes?
  - Segments
  - Pages (~4 KB on x86)
- Isolation?

EPIC
Institute of Technology
Powered by epam

Page Table

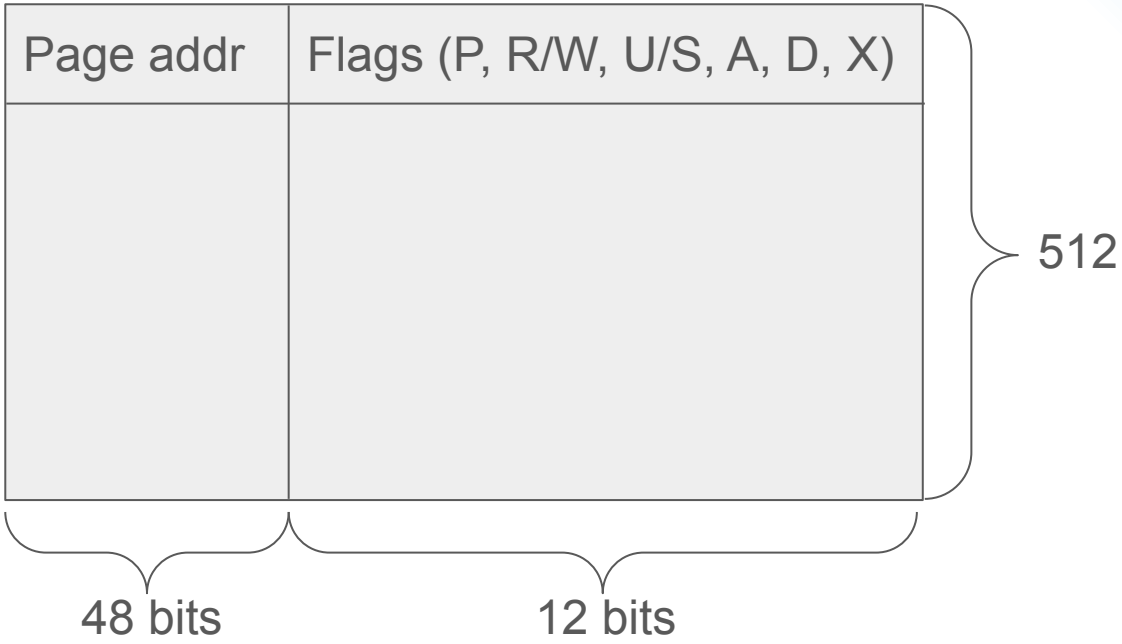

PTE Size = 1 page (4 KB)
Addresses up to 4KB*512 = 2MB

EPIC
Institute of Technology
Powered by epam

Page Table

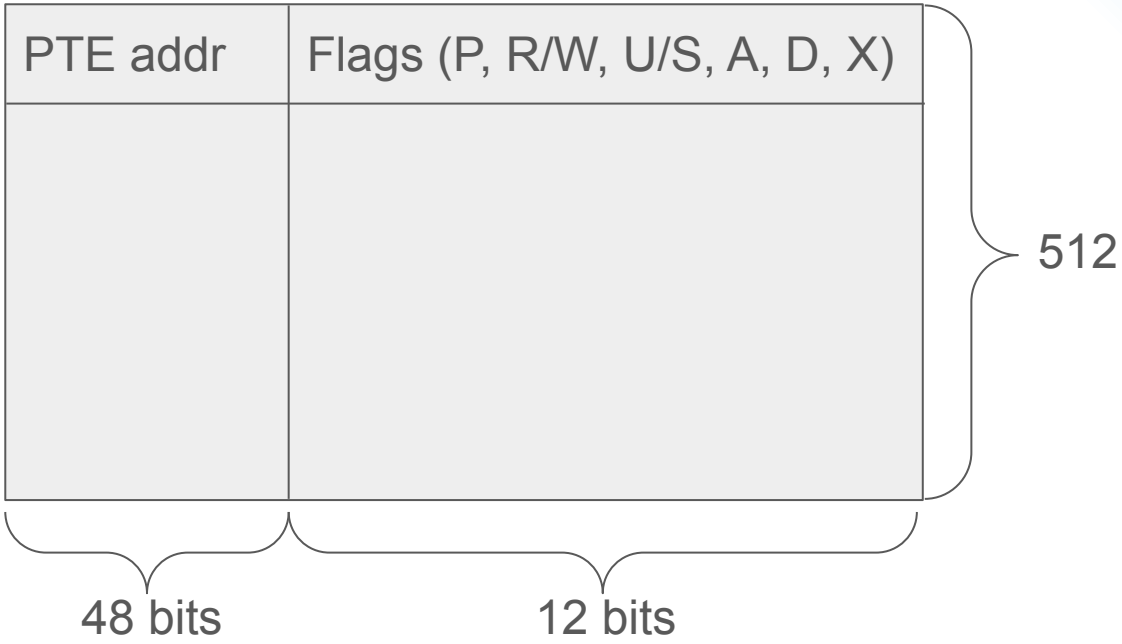

PDE Size = 1 page (4 KB)
Addresses up to 1GB

EPIC
Institute of Technology
Powered by epam

Page Table

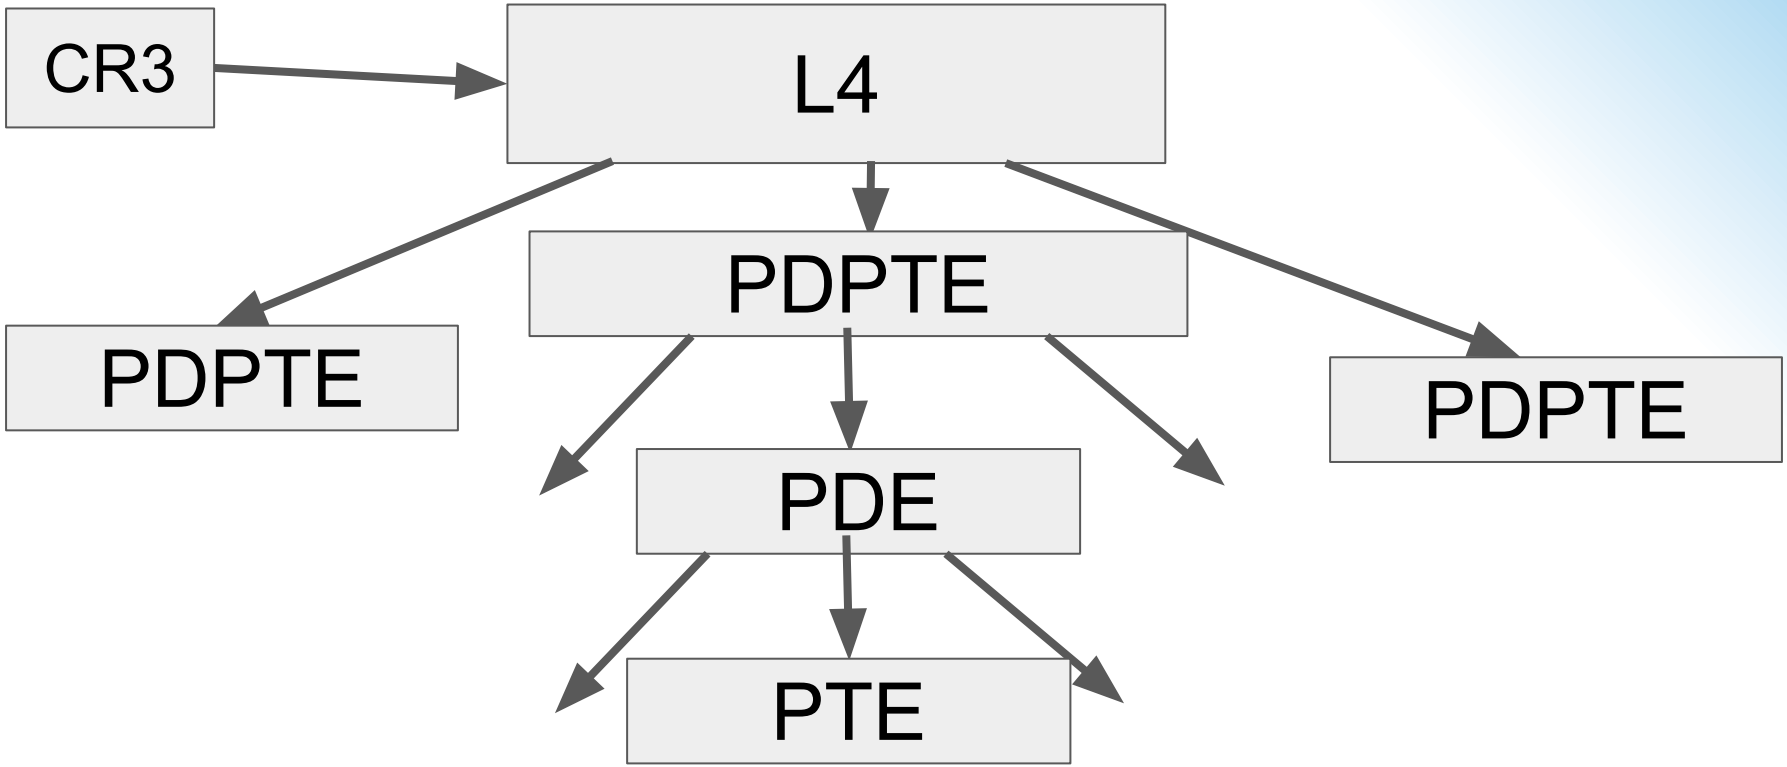

EPIC
Institute of Technology
Powered by epam

More Paging

- MMU
- Large pages (2 MB)
- Huge pages (1 GB)
- Caching (TLB)
- Total: 48 bit address

EPIC
Institute of Technology
Powered by epam

Isolation

- Each process gets its own page table
- Doesn’t see the memory of other processes
- Threads of one process get the same page table
- Shared memory: same physical page, different virtual

EPIC
Institute of Technology
Powered by <epam>

Fork, CoW

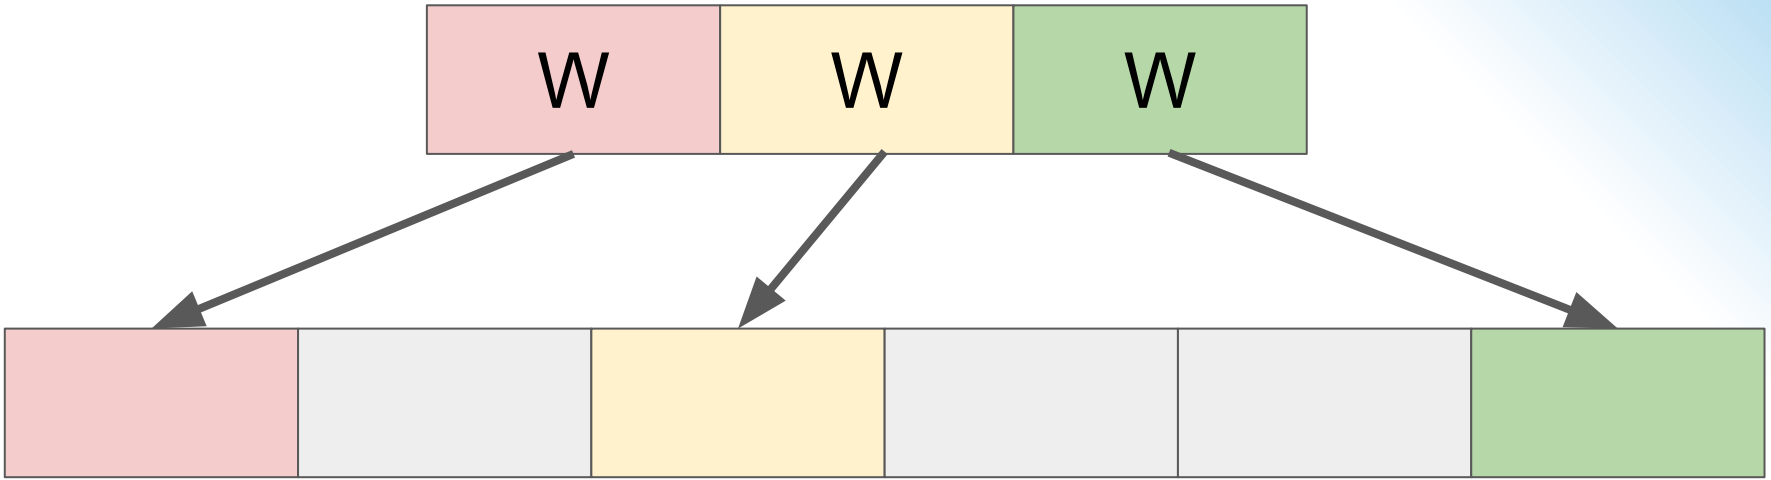

EPIC
Institute of Technology
Powered by epam

Fork, CoW

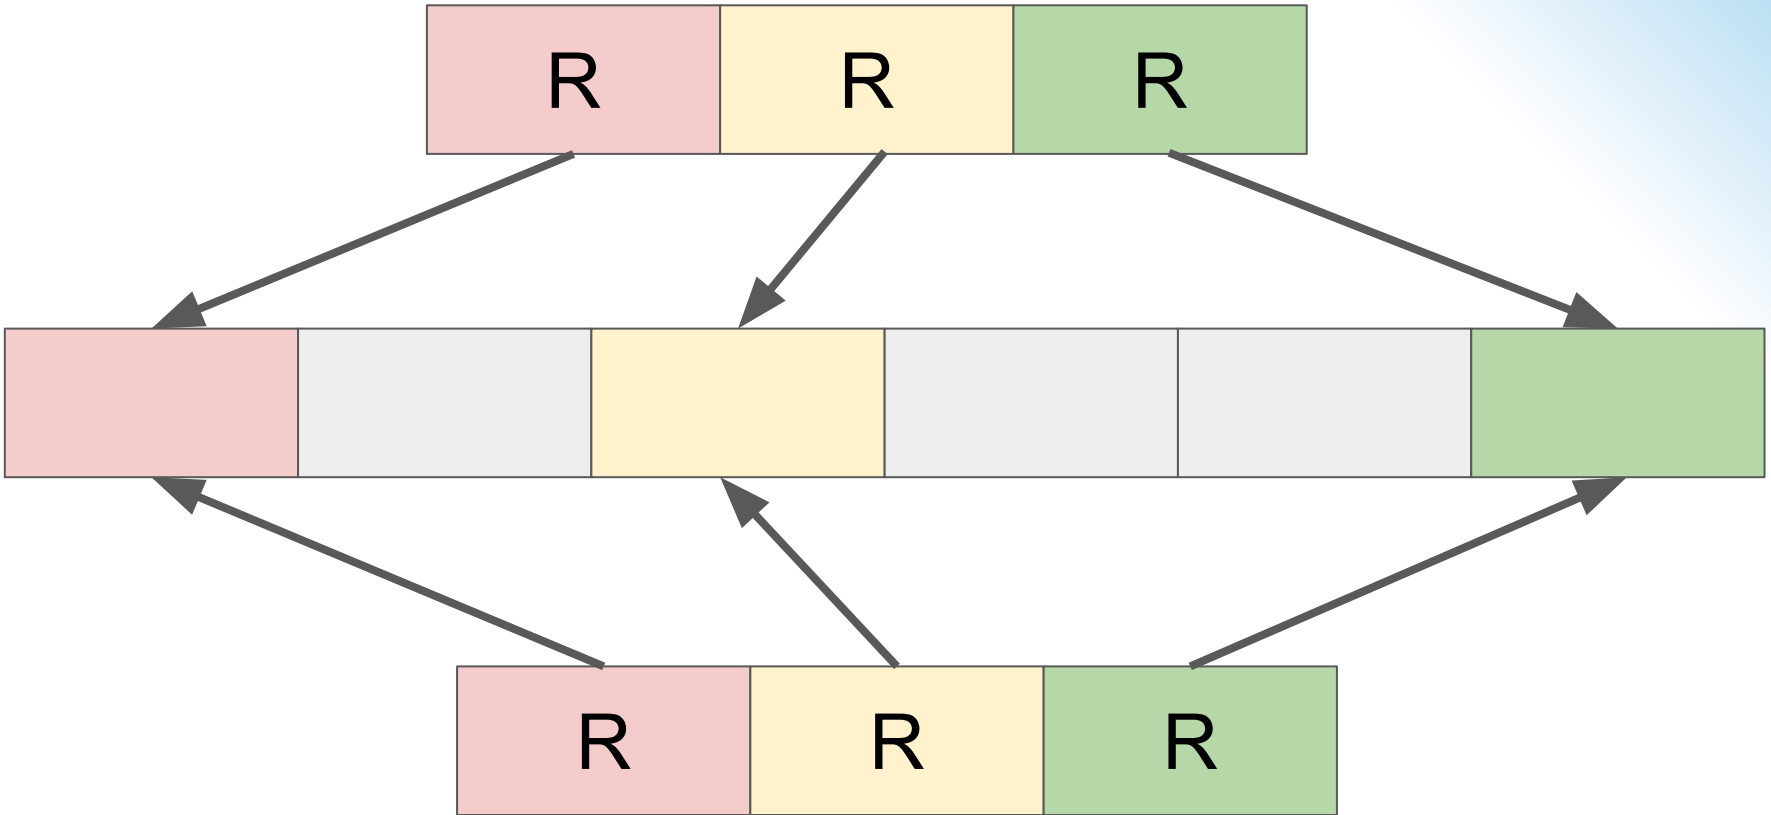

EPIC
Institute of Technology
Powered by epam

Fork, CoW

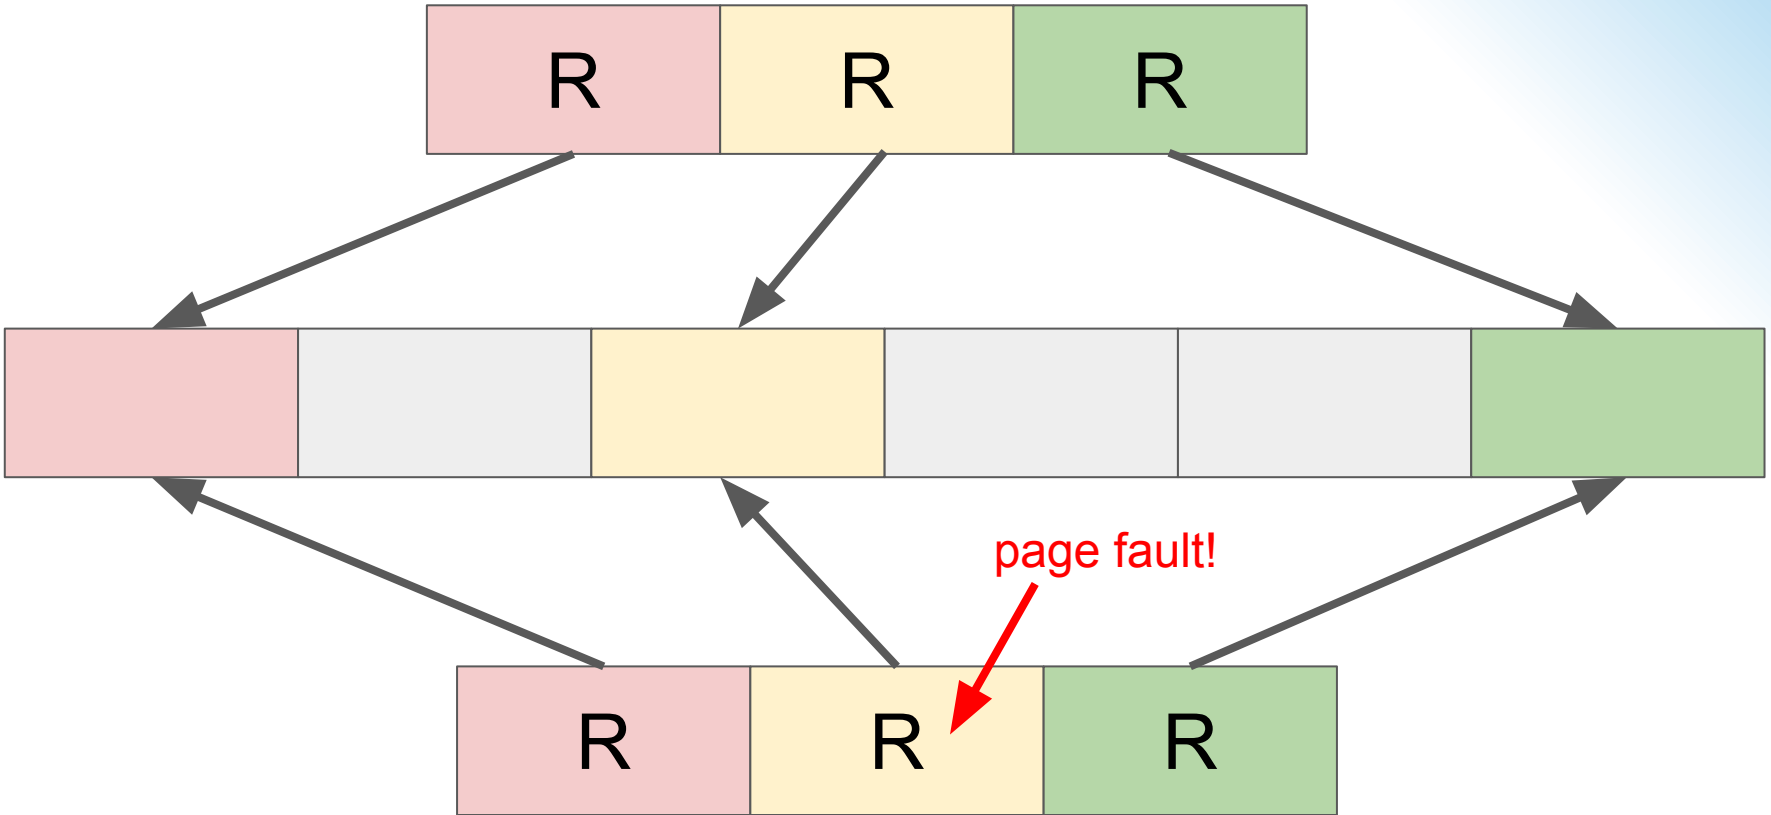

EPIC
Institute of Technology
Powered by epam

Fork, CoW

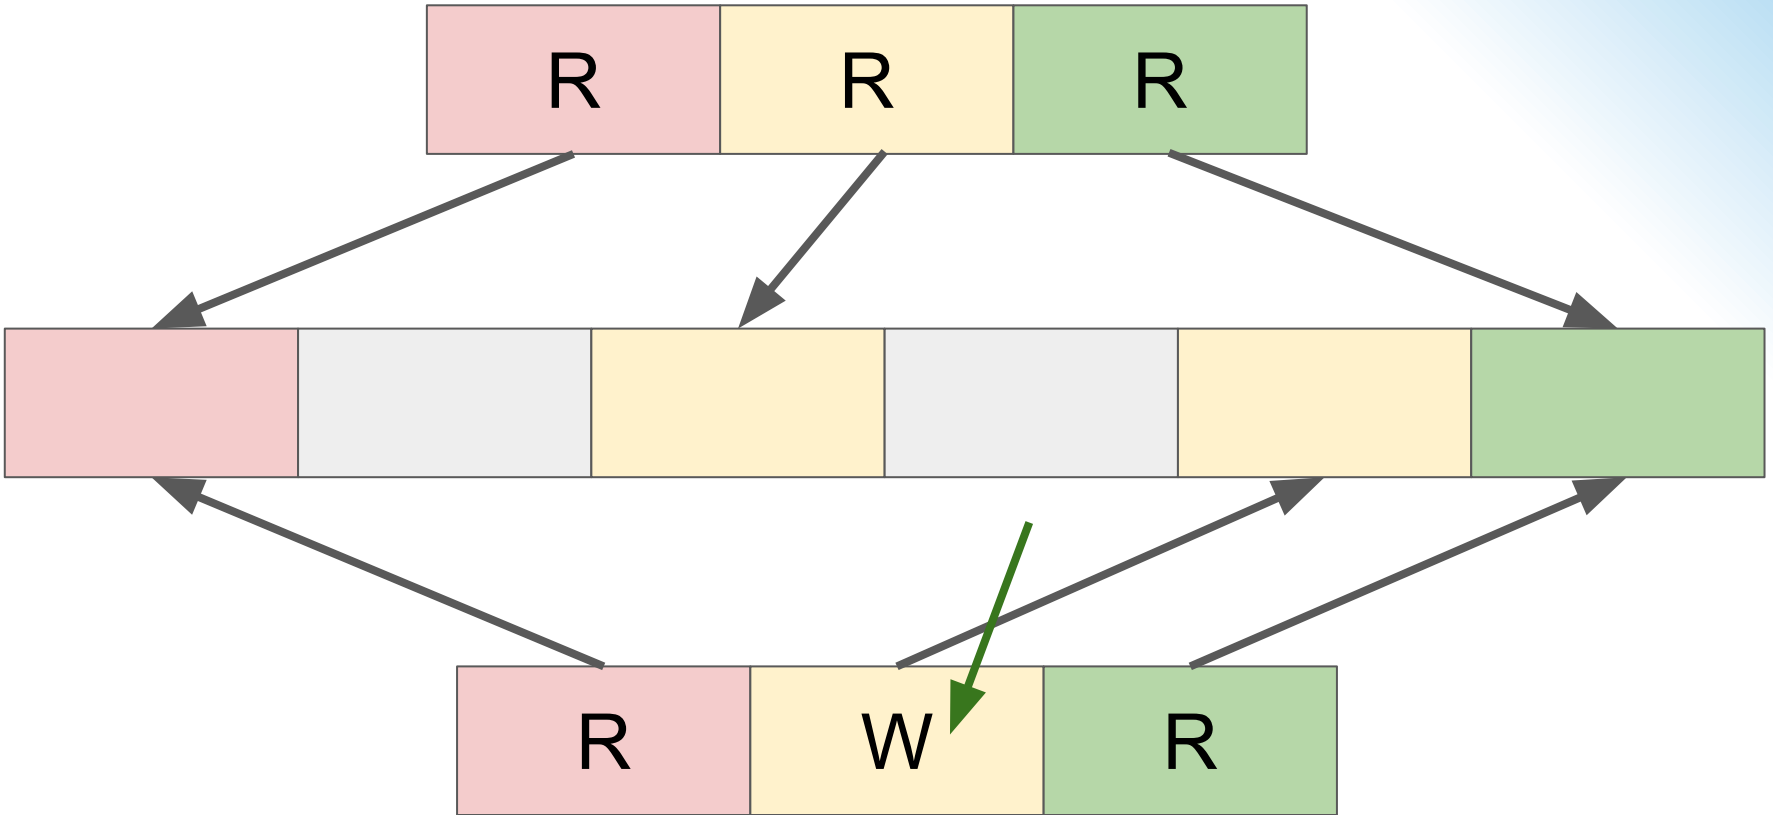

EPIC
Institute of Technology
Powered by <epam>

03

Memory Mapping

EPIC Institute of Technology
Powered by <epam>

Page cache

- Stores file contents in memory
- Speeds up access to data
- If the page is dirty → write back to disk

EPIC
Institute of Technology
Powered by epam

Memory Mapping

EPIC
Institute of Technology
Powered by <epam>

Memory Mapping

File on Disk

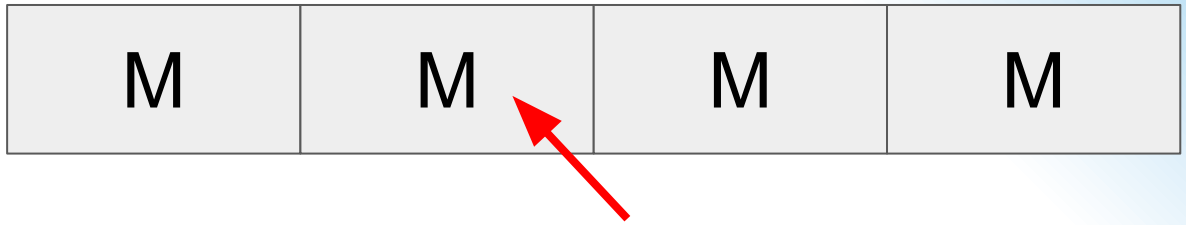

EPIC
Institute of Technology
Powered by <epam>

Memory Mapping

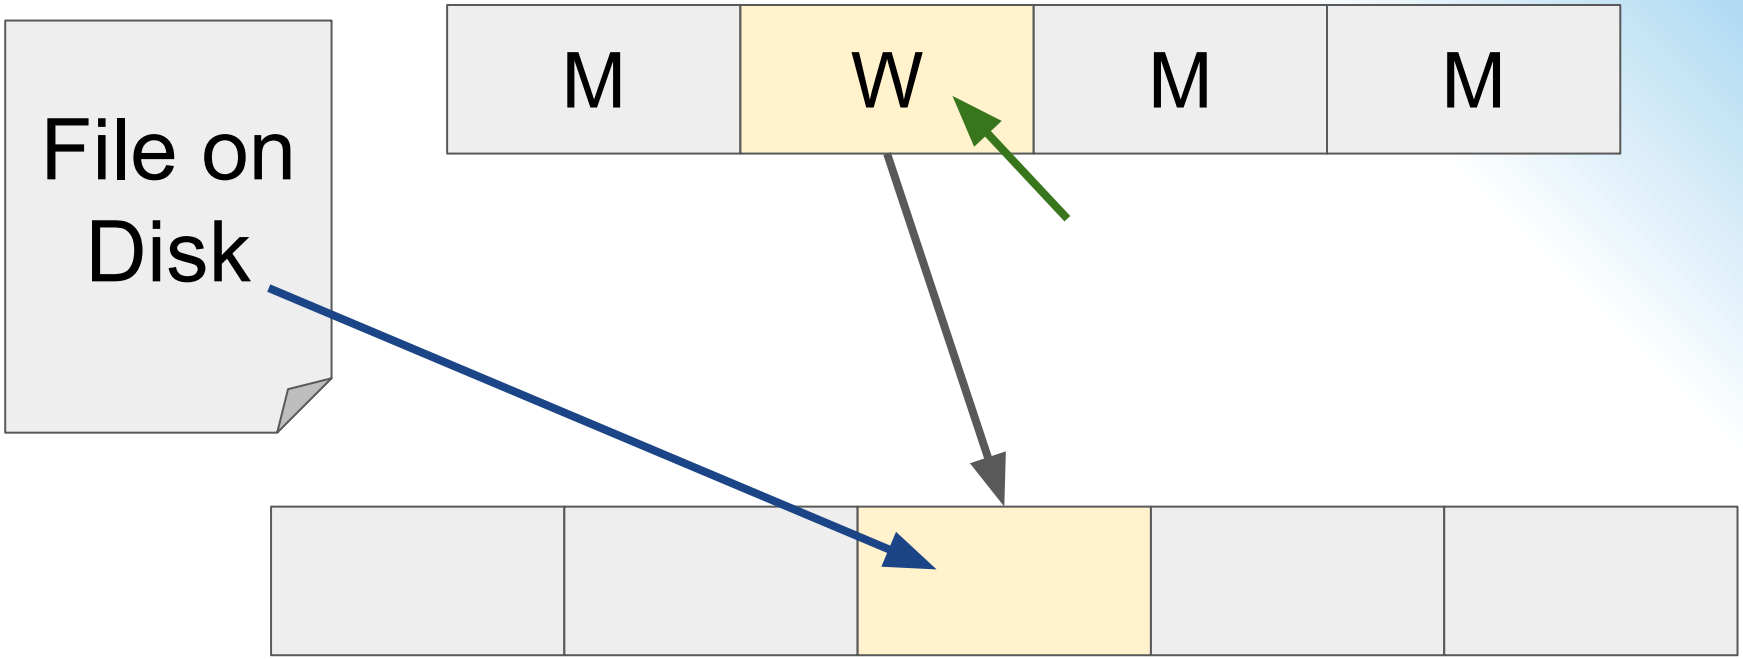

EPIC
Institute of Technology
Powered by epam

Memory Mapping

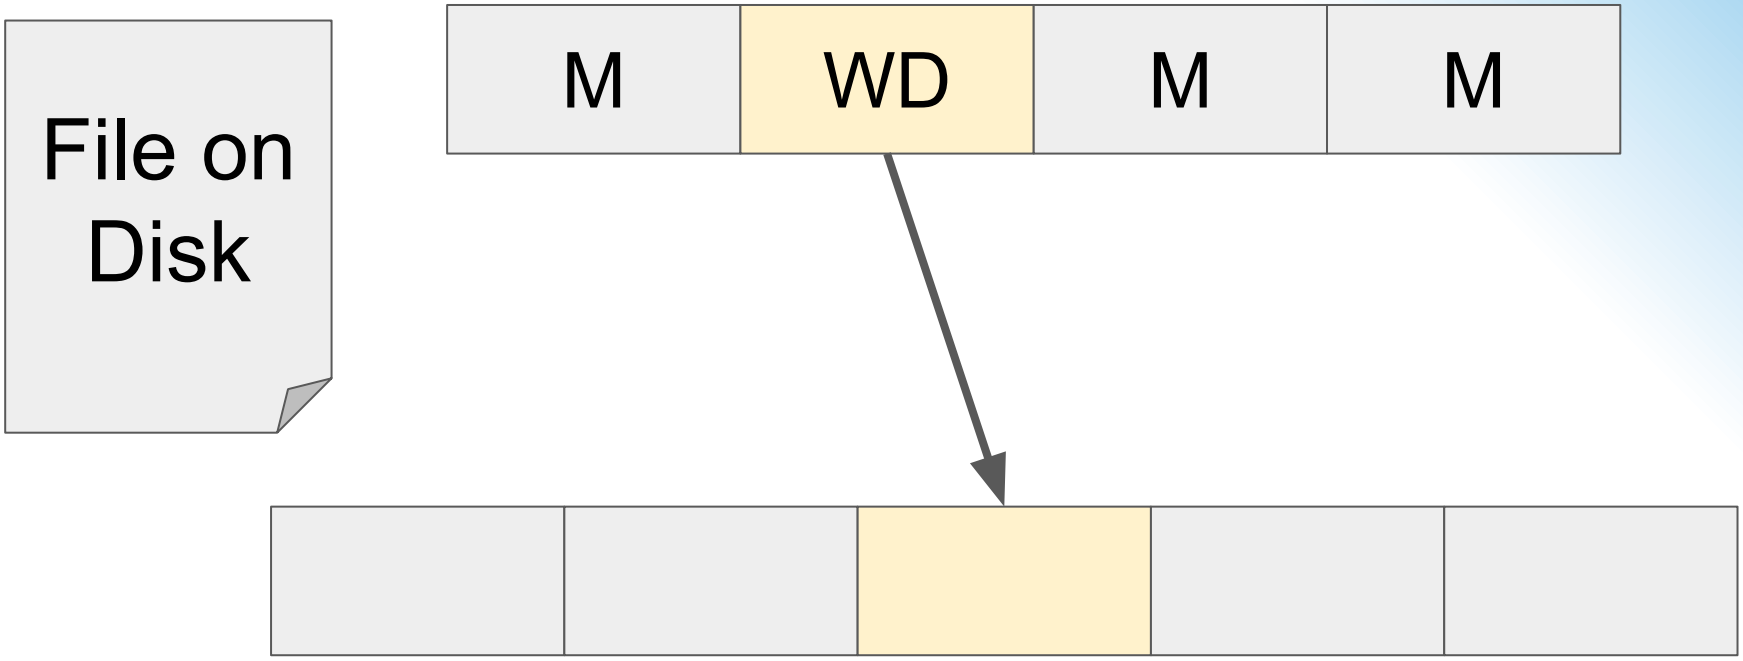

EPIC
Institute of Technology
Powered by <epam>

Memory Mapping