

# Introduction. CAP. Base

Course: Real-Time Backend

Lecturer: Gleb Lobanov

April 20, 2024

# Contents

**01 Distributed systems**

---

**02 Performance. Scalability**

---

**03 Fault tolerance and consistency**

---

**04 CAP theorem**

---

**05 Examples**

---

**06 Next steps**

---

01

# Distributed systems

In this section we will discuss what is “Distributed system” and why do we need it.

# What is a distributed system?

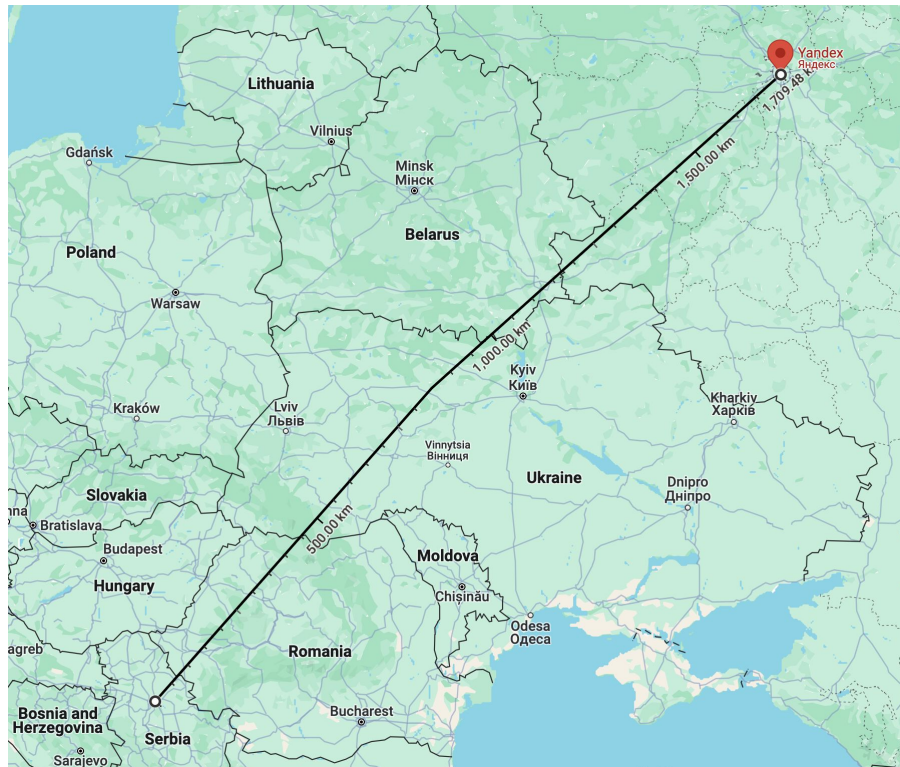
A distributed system is a system whose components are located on different networked computers, which communicate and coordinate their actions by passing messages to one another

*M. van Steen and A.S. Tanenbaum,  
Distributed Systems, 3rd ed.*

# Features

- Parallelism
  - More users / more data - solve tasks faster
- Fault tolerance
  - Continue working even a part of a system fails
- Physical distribution
  - 2 users in different countries
- Security / Isolation
  - Admin access/Regular access
- Separation
  - Premium users/Regular users

# Internet



299792 km/s - speed of light

200000 km/s - real speed

dist = 2000 km

$\text{min\_time} = 2000 / 200000 = 10 \text{ ms}$

$\text{both\_ways} = 10 * 2 = 20$

```
64 bytes from 188.35.21.10: icmp_seq=0 ttl=54 time=59.026 ms
64 bytes from 188.35.21.10: icmp_seq=1 ttl=54 time=57.408 ms
64 bytes from 188.35.21.10: icmp_seq=2 ttl=54 time=56.451 ms
64 bytes from 188.35.21.10: icmp_seq=3 ttl=54 time=56.620 ms
64 bytes from 188.35.21.10: icmp_seq=4 ttl=54 time=70.777 ms
64 bytes from 188.35.21.10: icmp_seq=5 ttl=54 time=58.574 ms
64 bytes from 188.35.21.10: icmp_seq=6 ttl=54 time=57.659 ms
```

# Challenges

- Resources
- Partial failures
- Performance

# Infrastructure

- Storage (GFS, S3, SQL db, NoSQL db, Clickhouse)
- Communication (Message brokers, GRPC, HTTP)
- Computation (Flink, MapReduce, Spark)



02

# Performance. Scalability

# Scalability

Task 1



Task 3

Task 2

# Scalability



# Scalability



Vertical scaling

# Scalability

Task 1

Task 2



Task 3

Task 4




Task ...




Horizontal  
scaling

2x machines = 2x times  
faster?

# Not all computations are independent





# or maybe even harder





Having storage also makes  
things more complicated







We also want to provide  
abstractions so  
client can work with distributed  
system conveniently



**EPIC**

Institute of Technology  
Powered by ePam





03

# Fault tolerance and consistency

# Fault tolerance

## What we mean by fault tolerance

- Availability
  - Under certain set of failures we still provide service
- Recoverability
  - Can recover from a failed state

## How to achieve

- Using permanent storage
- Replication



# Consistency


## Example: KV service

- Set(k, v)
- Get(k) -> v

# Consistency

## Example: KV service

- Set(k, v)
- Get(k) -> v




# Consistency

## Example: KV service

- Set(k, v)
- Get(k) -> v

k	v
abc	21


k	v
abc	20



# Consistency

## Example: KV service

- Set(k, v)
- Get(k) -> v



# Consistency

## Example: KV service

- Set(k, v)
- Get(k) -> v

k	v
abc	21

↕


k	v
abc	201

get(abc) → wait



**EPIC**Institute of Technology  
Powered by <epam>

# Example



04

# CAP theorem

# CAP theorem

A shared-data system can have at most two of the three following properties: **C**onsistency, **A**vailability, and tolerance to network **P**artitions



# On Consistency

Atomic, or **linearizable**, consistency is the condition expected by most web services today. Under this consistency guarantee, there must exist a total order on all operations such that each operation looks as if it were completed at a single instant. This is equivalent to requiring requests of the distributed shared memory to act as if they were executing on a single node, responding to operations one at a time.

# On Availability

For a distributed system to be continuously available, every request received by a non-failing node in the system must result in a response. That is, any algorithm used by the service must eventually terminate ... [When] qualified by the need for partition tolerance, this can be seen as a strong definition of availability: even when severe network failures occur, every request must terminate.

# On Partition tolerance



In order to model partition tolerance, the network will be allowed to lose arbitrarily many messages sent from one node to another. When a network is partitioned, all messages sent from nodes in one component of the partition to nodes in another component are lost.

You can't sacrifice  
partition tolerance!



05

# Examples




# No partition tolerance



# No partition tolerance






# No consistency





# No availability



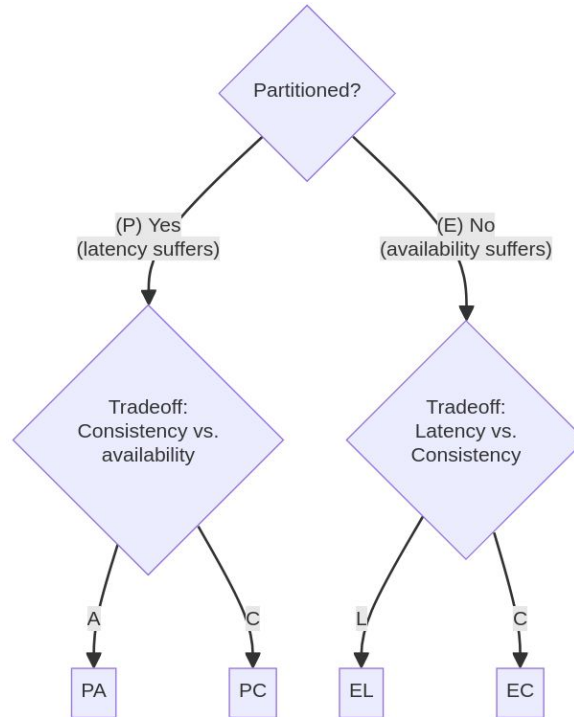
06

# What next?

# Next steps

- PACELC
- BASE
- ACID

# PACELC



if (P, then A or C, else L or C)

# That's All Folks!