

Intro to the subject

Course: Systems Architecture

Lecturer: Gleb Lobanov

February 19, 2024

Contents

01 **Introduction**

02 **Why O-notation is bad**

03 **AVX**

04 **Leaky abstraction**

01

Introduction

In this section, there will be a brief introduction about myself and a description of this course.

About me

Gleb Lobanov

Current Workplace – B2B YANGO

Tutoring – MIPT, Tinkoff, olympiad schools and so on

Tg/discord - glebodin,
if you need something special – feel free to ask

Grading

$$M_{\text{final}} = \min(5, \text{round}(0.7 \cdot M_{\text{practice}} + 0.3 \cdot M_{\text{individual_work}} + M_{\text{bonuses}}))$$

$$M_{\text{practice}} = \text{amount_solved}/\text{amount_necessary} * 5$$

$$M_{\text{individual_work}} = 0.4 \cdot M_{\text{midterm}} + 0.6 \cdot M_{\text{exam}}$$

M_{bonuses} may be given on various unpredictable special occasions. Almost never is negative 😊

Codestyle

Automatic:

- TAB is strictly prohibited
- Tabulation is 4 spaces
- Tabulation style is K&R, Linux Kernel

Manual:

- Meaningful naming
- Meaningful commenting
- One Code style for one file

Motivation

- Job in low-level applications(security, high load and so on)
- Less problems in every program

What pitfalls can you meet if you skip this course?

- Wrong understanding of asymptotics
- Bad memory usage in your programs (for example Cache Miss).
- Not knowing and not using SIMD, AVX and so on
- Leaky abstraction

02

Why O-notation is bad

In this section, we will explain why O-notation doesn't always apply in reality.

Problems of O-notation

- Different programming languages
- Different processor architecture
- Different cost of operations



EPIC

Institute of Technology
Powered by epam

Python

```
1. n = 128
2. d = [[0] * n] * n
3.
4. for i in range(n):
5.     for j in range(n):
6.         d[i][j] = i + j
7.
8. for k in range(n):
9.     for i in range(n):
10.        for j in range(n):
11.            d[i][j] = min(d[i][j], d[i][k] + d[k][j])
12.
```

Success #stdin #stdout 1.2s 9732KB

<https://ideone.com/3cU8kk>

C++

```
1. #include <iostream>
2.
3. using namespace std;
4. const int n = 1024;
5. int d[n][n];
6.
7. int main() {
8.
9.     for (int i = 0; i < n; i++) {
10.         for (int j = 0; j < n; j++) {
11.             d[i][j] = i + j;
12.         }
13.     }
14.
15.     for (int k = 0; k < n; k++) {
16.         for (int i = 0; i < n; i++) {
17.             for (int j = 0; j < n; j++) {
18.                 d[i][j] = min(d[i][j], d[i][k] + d[k][j]);
19.             }
20.         }
21.     }
22.
23. }
```

Success #stdin #stdout 1.28s 7672KB

<https://ideone.com/2yAssk>



EPIC

Institute of Technology
Powered by epam

Rust

```
1.  use std::cmp;
2.
3.  fn main() {
4.      const n: usize = 1024;
5.      let mut d = [[0; n]; n];
6.      for i in 0 .. n - 1 {
7.          for j in 0 .. n - 1 {
8.              d[i][j] = i + j
9.          }
10.     }
11.     for k in 0 .. n - 1 {
12.         for i in 0 .. n - 1 {
13.             for j in 0 .. n - 1 {
14.                 d[i][j] = cmp::min(d[i][j], d[i][k] + d[k][j])
15.             }
16.         }
17.     }
18. }
```

Success #stdin #stdout 1.13s 10064KB

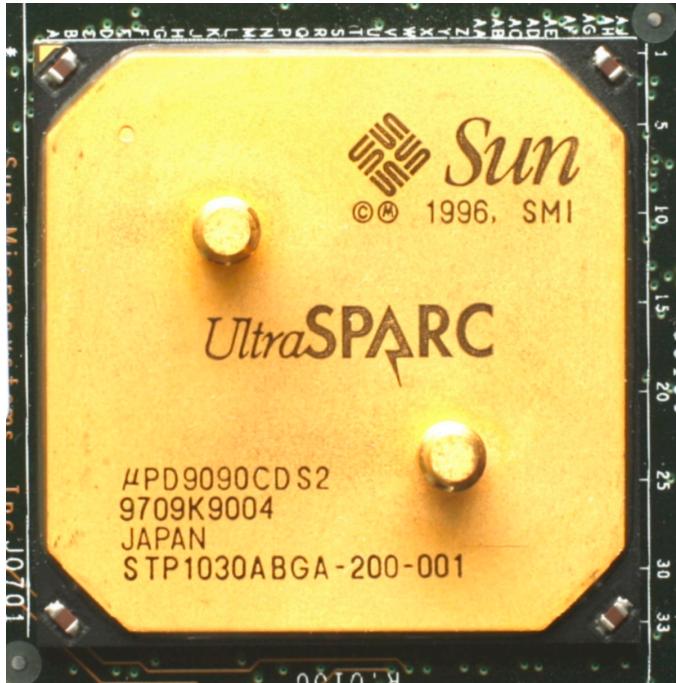
stdin

Standard input is empty

stdout

Standard output is empty

Computer architecture



- RISC
- CISC
- MISC

Tricky example

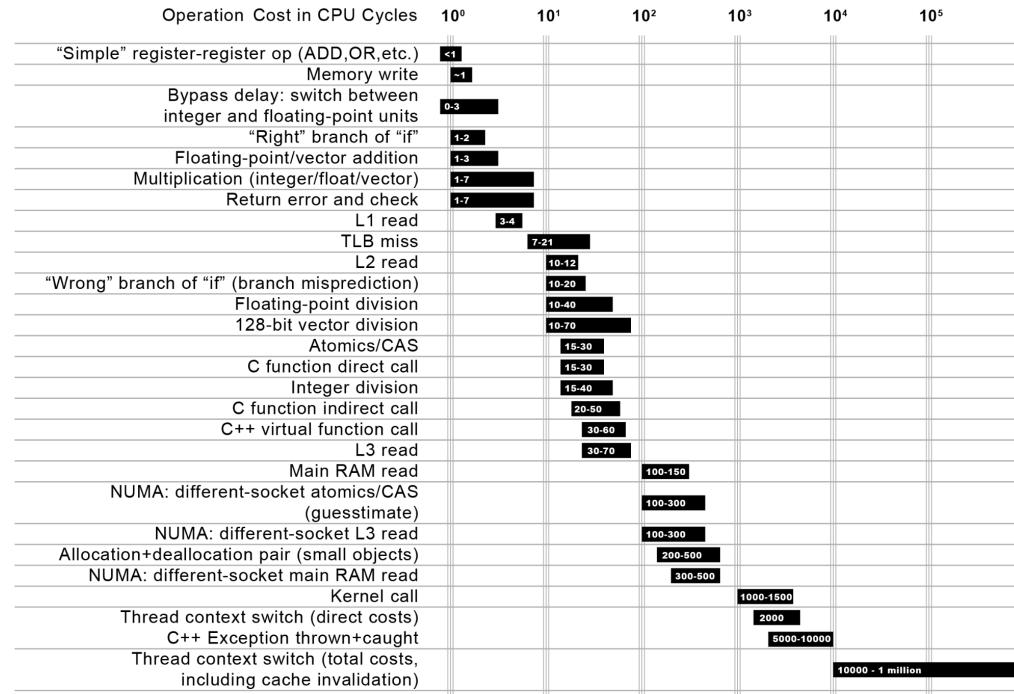
- Intel
compare_exchange
 - ARM
LL/SC

<https://en.wikipedia.org/wiki/Load-link/store-conditional>
 - ARM
LL/SC
- <https://en.wikipedia.org/wiki/Load-link/store-conditional>

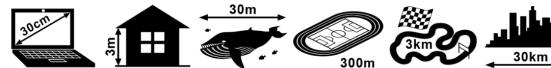
Elementary operations in O-notation

- Bit operations (bitwise and, bitwise or, ...);
- arithmetic operations (add, multiply, divide, ...);
- some comparison operations (compare, greater, less, ..);
- some logical operations (and, or, ...);
- some memory operations (take from memory, put to memory, ...);
- some IO operations (read, write, ...);
- If, exception and more other.

Not all CPU operations are created equal



Distance which light travels while the operation is performed



Alternatives



PRAM



External Memory Model

HDD vs SSD

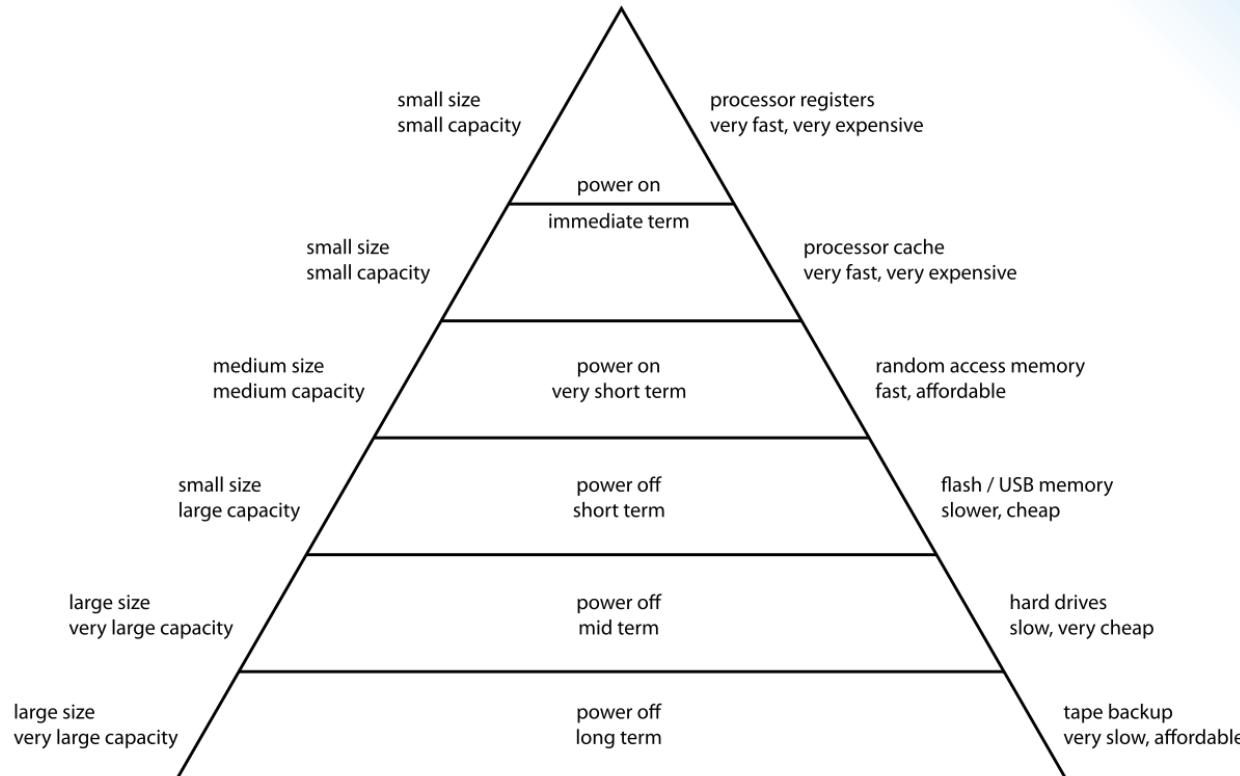


HDD



SSD

Computer Memory Hierarchy



Real Life

Here is an approximate comparison table for commodity hardware in 2021

Type	<i>M</i>	<i>B</i>	Latency	Bandwidth	\$/GB/mo ¹
L1	10K	64B	2ns	80G/s	-
L2	100K	64B	5ns	40G/s	-
L3	1M/core	64B	20ns	20G/s	-
RAM	GBs	64B	100ns	10G/s	1.5
SSD	TBs	4K	0.1ms	5G/s	0.17
HDD	TBs	-	10ms	1G/s	0.04
S3	∞	-	150ms	∞	0.02 ²

Caches as beer



Holly Cummins (holly_cummins@hachyderm.io)
[@holly_cummins](https://twitter.com/holly_cummins) · [Follow](#)

L1 cache is a beer in hand, L3 is fridge, main memory is walking to the store, disk access is flying to another country for beer. [@netOpyr](#)

03

AVX

In this section, we will discuss vectorized computations.

$\text{type}(ai) = \text{type}(bi) = \text{double}$

Regular

A	a1	a2	a3	a4	a5	a6	a7	a8
---	----	----	----	----	----	----	----	----

+

B	b1	b2	b3	b4	b5	b6	b7	b8
---	----	----	----	----	----	----	----	----

=

A + B	a1+b1	a2+b2	a3+b3	a4+b4	a5+b5	a6+b6	a7+b7	a8+b8
-------	-------	-------	-------	-------	-------	-------	-------	-------



EPIC

Institute of Technology
Powered by epam

type(ai) = type(bi) = double

SSE

A

a1	a2	a3	a4	a5	a6	a7	a8
----	----	----	----	----	----	----	----

+

B

b1	b2	b3	b4	b5	b6	b7	b8
----	----	----	----	----	----	----	----

=

A + B

a1+b1	a2+b2	a3+b3	a4+b4	a5+b5	a6+b6	a7+b7	a8+b8
-------	-------	-------	-------	-------	-------	-------	-------



EPIC

Institute of Technology
Powered by epam

type(ai) = type(bi) = double

AVX

A

a1	a2	a3	a4	a5	a6	a7	a8
----	----	----	----	----	----	----	----

+

B

b1	b2	b3	b4	b5	b6	b7	b8
----	----	----	----	----	----	----	----

=

A + B

a1+b1	a2+b2	a3+b3	a4+b4	a5+b5	a6+b6	a7+b7	a8+b8
-------	-------	-------	-------	-------	-------	-------	-------



EPIC

Institute of Technology
Powered by epam

type(ai) = type(bi) = double

AVX512

A	a1	a2	a3	a4	a5	a6	a7	a8
+								
B	b1	b2	b3	b4	b5	b6	b7	b8
=								
A + B	a1+b1	a2+b2	a3+b3	a4+b4	a5+b5	a6+b6	a7+b7	a8+b8

Instruction Set

- MMX
- SSE family
- AVX family
- AVX-512 family
- AMX family
- SVM
- Other

Categories

- Application-Targeted
- Arithmetic
- Bit Manipulation
- Cast
- Compare
- Convert
- Cryptography
- Elementary Math Functions
- General Support
- Load
- Logical
- Mask
- Miscellaneous
- Move
- OS-Targeted
- Probability/Statistics
- Random

 Search Intel Intrinsics

[source link](#)

```
void _mm_2intersect_epi32 (_m128i a, _m128i b, _mmask8* k1, _mmask8* k2) vp2intersectd
void _mm256_2intersect_epi32 (_m256i a, _m256i b, _mmask8* k1, _mmask8* k2) vp2intersectd
void _mm512_2intersect_epi32 (_m512i a, _m512i b, _mmask16* k1, _mmask16* k2) vp2intersectd
void _mm_2intersect_epi64 (_m128i a, _m128i b, _mmask8* k1, _mmask8* k2) vp2intersectq
void _mm256_2intersect_epi64 (_m256i a, _m256i b, _mmask8* k1, _mmask8* k2) vp2intersectq
void _mm512_2intersect_epi64 (_m512i a, _m512i b, _mmask8* k1, _mmask8* k2) vp2intersectq
void _aadd_i32 (int* __A, int __B) aadd
void _aadd_i64 (__int64* __A, __int64 __B) aadd
void _aand_i32 (int* __A, int __B) aand
void _aand_i64 (__int64* __A, __int64 __B) aand
__m128i _mm_abs_epi16 (__m128i a) pabsw
__m128i _mm_mask_abs_epi16 (__m128i src, _mmask8 k, __m128i a) vpabsw
__m128i _mm_maskz_abs_epi16 (_mmask8 k, __m128i a) vpabsw
__m256i _mm256_abs_epi16 (__m256i a) vpabsw
__m256i _mm256_mask_abs_epi16 (__m256i src, _mmask16 k, __m256i a) vpabsw
__m256i _mm256_maskz_abs_epi16 (_mmask16 k, __m256i a) vpabsw
__m512i _mm512_abs_epi16 (__m512i a) vpabsw
__m512i _mm512_mask_abs_epi16 (__m512i src, _mmask32 k, __m512i a) vpabsw
__m512i _mm512_maskz_abs_epi16 (_mmask32 k, __m512i a) vpabsw
__m128i _mm_abs_epi32 (__m128i a) pabsd
```

Without AVX

```
1. #include <bits/stdc++.h>
2.
3. using namespace std;
4. const int SIZE = 50000000;
5.
6. double a[SIZE], b[SIZE], c[SIZE];
7.
8. int main() {
9.     for (int i = 0; i < SIZE; i++) {
10.         a[i] = b[i] = i;
11.     }
12.     for (int i = 0; i < SIZE; i++) {
13.         c[i] = a[i] + b[i];
14.     }
15.
16.     return 0;
17. }
```

Success #stdin #stdout 0.27s 1175544KB

 stdin

Standard input is empty

 stdout

Standard output is empty

With AVX

```
1. #pragma GCC target("avx2")
2. #include <bits/stdc++.h>
3. #include <immintrin.h>
4.
5. using namespace std;
6. const int SIZE = 50000000;
7.
8. double a[SIZE], b[SIZE], c[SIZE];
9.
10. int main() {
11.     for (int i = 0; i < SIZE; i += 4) {
12.         a[i] = b[i] = i;
13.     }
14.     for (int i = 0; i < SIZE; i += 4) {
15.         __m256d x = _mm256_loadu_pd(&a[i]);
16.         __m256d y = _mm256_loadu_pd(&b[i]);
17.         __m256d z = _mm256_add_pd(x, y);
18.         _mm256_storeu_pd(&c[i], z);
19.     }
20.
21.     return 0;
22. }
```

Real Task

Let's say we have a task:

- 1) Find the number of numbers $< x$ on the segment
- 2) Change the number.

In our algorithms course, we've already learned how to do this quite efficiently (using segment tree, Cartesian tree, or a sqrt decomposition).

But what if there's a much simpler way to do it?



Real Task

```
if (t == 1) {
    a--;
    int l = (a / d) * d;
    int r = min(n, l + d);
    int pb = lower_bound(g + l, g + r, s[a]) - g;
    if (s[a] < b) {
        for (; pb < r - 1 && g[pb + 1] < b; pb++)
            g[pb] = g[pb + 1];
    }
    g[pb] = b;
} else {
    for (; pb > l && g[pb - 1] > b; pb--) {
        g[pb] = g[pb - 1];
    }
    g[pb] = b;
}
s[a] = b;
} else {
    cin >> c;
    b--;
    int l = (b / d) * d;
    int r = min(n, l + d);
    int ans = 0;
    if (b > ((l + r) >> 1) || r >= c) {
        int t = min(r, c);
        for (int i = b; i < t; i++) {
            if (s[i] < a) {f
                ans--;
            }
        }
    }
}
```



EPIC

Institute of Technology
Powered by epam

Real Task

```
    cin >> tp;
    if (tp == 1) {
        int i, x;
        cin >> i >> x;
        a[i - 1] = x;
    } else {
        int m, l, r;
        cin >> m >> l >> r;
        int cl = 0;
        for (int i = l - 1; i < r; ++i)
            cl += a[i] < m;
```



EPIC

Institute of Technology
Powered by epam

Comparison

[ANSWER](#)

nsqrt_opt.cpp	stupid_solution.cpp
<u>OK</u> 0 / 1	<u>OK</u> 0 / 0
<u>OK</u> 0 / 1 / 1.00000	<u>OK</u> 0 / 0 / 1.00000
<u>OK</u> 15 / 1 / 1.00000	<u>OK</u> 0 / 0 / 1.00000
<u>OK</u> 15 / 1 / 1.00000	<u>OK</u> 15 / 0 / 1.00000
<u>OK</u> 467 / 1 / 1.00000	<u>OK</u> 3946 / 0 / 1.00000
<u>OK</u> 1122 / 1 / 1.00000	<u>OK</u> 7675 / 0 / 1.00000
<u>OK</u> 1107 / 1 / 1.00000	<u>TL</u> 15000 / 0 / 1.00000
<u>OK</u> 763 / 1 / 1.00000	<u>OK</u> 4055 / 0 / 1.00000
<u>OK</u> 15 / 1 / 1.00000	<u>OK</u> 0 / 0 / 1.00000
<u>OK</u> 0 / 1 / 1.00000	<u>OK</u> 0 / 0 / 1.00000
<u>OK</u> 0 / 1 / 1.00000	<u>OK</u> 15 / 0 / 1.00000
<u>OK</u> 483 / 1 / 1.00000	<u>OK</u> 2901 / 0 / 1.00000
<u>OK</u> 1076 / 1 / 1.00000	<u>TL</u> 15000 / 0 / 1.00000
<u>OK</u> 1075 / 1 / 1.00000	<u>TS</u> 0 / 0 / 1.00000
<u>OK</u> 0 / 1 / 1.00000	<u>OK</u> 0 / 0 / 1.00000
<u>OK</u> 1091 / 1 / 1.00000	<u>TS</u> 0 / 0 / 1.00000
<u>OK</u> 46 / 1 / 1.00000	<u>OK</u> 187 / 0 / 1.00000
<u>OK</u> 15 / 1 / 1.00000	<u>OK</u> 78 / 0 / 1.00000
<u>OK</u> 1060 / 1 / 1.00000	<u>TS</u> 0 / 0 / 1.00000
<u>OK</u> 1076 / 1 / 1.00000	<u>TL</u> 15000 / 0 / 1.00000
<u>OK</u> 171 / 1 / 1.00000	<u>OK</u> 919 / 0 / 1.00000



EPIC

Institute of Technology
Powered by epam

Eazy AVX

```
1 #pragma GCC target("avx2")
2 #pragma GCC optimize("O3")
3
```



EPIC

Institute of Technology
Powered by epam

nsqrt_opt.cpp	pragma_avx.cpp	stupid_solution.cpp
<u>OK</u> 0 / 1	<u>OK</u> 0 / 0	<u>OK</u> 0 / 0
<u>OK</u> 0 / 1 / 1.00000	<u>OK</u> 0 / 0 / 1.00000	<u>OK</u> 0 / 0 / 1.00000
<u>OK</u> 15 / 1 / 1.00000	<u>OK</u> 0 / 0 / 1.00000	<u>OK</u> 0 / 0 / 1.00000
<u>OK</u> 15 / 1 / 1.00000	<u>OK</u> 15 / 0 / 1.00000	<u>OK</u> 15 / 0 / 1.00000
<u>OK</u> 467 / 1 / 1.00000	<u>OK</u> 1029 / 0 / 1.00000	<u>OK</u> 3946 / 0 / 1.00000
<u>OK</u> 1122 / 1 / 1.00000	<u>OK</u> 1949 / 0 / 1.00000	<u>OK</u> 7675 / 0 / 1.00000
<u>OK</u> 1107 / 1 / 1.00000	<u>OK</u> 5444 / 0 / 1.00000	<u>TL</u> 15000 / 0 / 1.00000
<u>OK</u> 763 / 1 / 1.00000	<u>OK</u> 1091 / 0 / 1.00000	<u>OK</u> 4055 / 0 / 1.00000
<u>OK</u> 15 / 1 / 1.00000	<u>OK</u> 0 / 0 / 1.00000	<u>OK</u> 0 / 0 / 1.00000
<u>OK</u> 0 / 1 / 1.00000	<u>OK</u> 0 / 0 / 1.00000	<u>OK</u> 0 / 0 / 1.00000
<u>OK</u> 0 / 1 / 1.00000	<u>OK</u> 0 / 0 / 1.00000	<u>OK</u> 15 / 0 / 1.00000
<u>OK</u> 483 / 1 / 1.00000	<u>OK</u> 732 / 0 / 1.00000	<u>OK</u> 2901 / 0 / 1.00000
<u>OK</u> 1076 / 1 / 1.00000	<u>OK</u> 5287 / 0 / 1.00000	<u>TL</u> 15000 / 0 / 1.00000
<u>OK</u> 1075 / 1 / 1.00000	<u>OK</u> 5444 / 0 / 1.00000	<u>TL</u> 15000 / 0 / 1.00000
<u>OK</u> 0 / 1 / 1.00000	<u>OK</u> 0 / 0 / 1.00000	<u>OK</u> 0 / 0 / 1.00000
<u>OK</u> 1091 / 1 / 1.00000	<u>OK</u> 5101 / 0 / 1.00000	<u>TL</u> 15000 / 0 / 1.00000
<u>OK</u> 46 / 1 / 1.00000	<u>OK</u> 61 / 0 / 1.00000	<u>OK</u> 187 / 0 / 1.00000
<u>OK</u> 15 / 1 / 1.00000	<u>OK</u> 30 / 0 / 1.00000	<u>OK</u> 78 / 0 / 1.00000
<u>OK</u> 1060 / 1 / 1.00000	<u>OK</u> 5101 / 0 / 1.00000	<u>TL</u> 15000 / 0 / 1.00000
<u>OK</u> 1076 / 1 / 1.00000	<u>OK</u> 4976 / 0 / 1.00000	<u>TL</u> 15000 / 0 / 1.00000
<u>OK</u> 171 / 1 / 1.00000	<u>OK</u> 249 / 0 / 1.00000	<u>OK</u> 919 / 0 / 1.00000

AVX

```
j
__m256i mm = _mm256_set1_epi32(m);
__m256i ansm = _mm256_set1_epi32(0);
for (int i = l; i + B <= r; i += B) {
    __m256i b0 = _mm256_lddqu_si256(reinterpret_cast<__m256i *>(a + i));
    __m256i b1 = _mm256_lddqu_si256(reinterpret_cast<__m256i *>(a + i + 8));
    b0 = _mm256_cmpgt_epi32(mm, b0);
    b1 = _mm256_cmpgt_epi32(mm, b1);
    ansm = _mm256_sub_epi32(ansm, b0);
    ansm = _mm256_sub_epi32(ansm, b1);
}
ans += _mm256_extract_epi32(ansm, 0);
ans += _mm256_extract_epi32(ansm, 1);
ans += _mm256_extract_epi32(ansm, 2);
ans += _mm256_extract_epi32(ansm, 3);
ans += _mm256_extract_epi32(ansm, 4);
ans += _mm256_extract_epi32(ansm, 5);
ans += _mm256_extract_epi32(ansm, 6);
ans += _mm256_extract_epi32(ansm, 7);
```

Extract

```
_int32 _mm256_extract_epi32 (_m256i a, const int index)
```

Synopsis

```
_int32 _mm256_extract_epi32 (_m256i a, const int index)
#include <immintrin.h>
```

Instruction: **Sequence**

CPUID Flags: AVX

Description

Extract a 32-bit integer from `a`, selected with `index`, and store the result in `dst`.

Operation

```
dst[31:0] := (a[255:0] >> (index[2:0] * 32))[31:0]
```

Set

```
_m256i _mm256_set1_epi32 (int a)
```

Synopsis

```
_m256i _mm256_set1_epi32 (int a)
#include <immintrin.h>
```

Instruction: **Sequence**

CPUID Flags: AVX

Description

Broadcast 32-bit integer `a` to all elements of `dst`. This intrinsic may generate the `vpbroadcastd`.

Operation

```
FOR j := 0 to 7
    i := j*32
    dst[i+31:i] := a[31:0]
ENDFOR
dst[MAX:256] := 0
```

Load

```
__m256i _mm256_lddqu_si256 (__m256i const * mem_addr)
```

vllddqu

Synopsis

```
__m256i _mm256_lddqu_si256 (__m256i const * mem_addr)
#include <immintrin.h>
Instruction: vllddqu ymm, m256
CPUID Flags: AVX
```

Description

Load 256-bits of integer data from unaligned memory into `dst`. This intrinsic may perform better than `_mm256_loadu_si256` when the data crosses a cache line boundary.

Operation

```
dst[255:0] := MEM[mem_addr+255:mem_addr]
dst[MAX:256] := 0
```

Compare

```
__m256i __mm256_cmpgt_epi32 (__m256i a, __m256i b)
```

Synopsis

```
__m256i __mm256_cmpgt_epi32 (__m256i a, __m256i b)
```

```
#include <immintrin.h>
```

```
Instruction: vpcmpqtd ymm, ymm, ymm
```

```
CPUID Flags: AVX2
```

Description

Compare packed signed 32-bit integers in `a` and `b` for greater-than, and store the results in `dst`.

Operation

```
FOR j := 0 to 7
    i := j*32
    dst[i+31:i] := ( a[i+31:i] > b[i+31:i] ) ? 0xFFFFFFFF : 0
ENDFOR
dst[MAX:256] := 0
```



EPIC

Institute of Technology
Powered by epam

AVX

honest_avx_opt.cpp	nsqrt_opt.cpp
<u>OK</u> 15 / 1	<u>OK</u> 0 / 1
<u>OK</u> 0 / 1 / 1.00000	<u>OK</u> 0 / 1 / 1.00000
<u>OK</u> 0 / 1 / 1.00000	<u>OK</u> 15 / 1 / 1.00000
<u>OK</u> 15 / 1 / 1.00000	<u>OK</u> 15 / 1 / 1.00000
<u>OK</u> 654 / 5 / 1.00000	<u>OK</u> 467 / 1 / 1.00000
<u>OK</u> 904 / 5 / 1.00000	<u>OK</u> 1122 / 1 / 1.00000
<u>OK</u> 1622 / 5 / 1.00000	<u>OK</u> 1107 / 1 / 1.00000
<u>OK</u> 717 / 5 / 1.00000	<u>OK</u> 763 / 1 / 1.00000
<u>OK</u> 0 / 1 / 1.00000	<u>OK</u> 15 / 1 / 1.00000
<u>OK</u> 0 / 1 / 1.00000	<u>OK</u> 0 / 1 / 1.00000
<u>OK</u> 0 / 1 / 1.00000	<u>OK</u> 0 / 1 / 1.00000
<u>OK</u> 452 / 5 / 1.00000	<u>OK</u> 483 / 1 / 1.00000
<u>OK</u> 1621 / 5 / 1.00000	<u>OK</u> 1076 / 1 / 1.00000
<u>OK</u> 1590 / 5 / 1.00000	<u>OK</u> 1075 / 1 / 1.00000
<u>OK</u> 0 / 1 / 1.00000	<u>OK</u> 0 / 1 / 1.00000
<u>OK</u> 1590 / 5 / 1.00000	<u>OK</u> 1091 / 1 / 1.00000
<u>OK</u> 62 / 1 / 1.00000	<u>OK</u> 46 / 1 / 1.00000
<u>OK</u> 15 / 1 / 1.00000	<u>OK</u> 15 / 1 / 1.00000
<u>OK</u> 1575 / 5 / 1.00000	<u>OK</u> 1060 / 1 / 1.00000
<u>OK</u> 1544 / 5 / 1.00000	<u>OK</u> 1076 / 1 / 1.00000
<u>OK</u> 171 / 2 / 1.00000	<u>OK</u> 171 / 1 / 1.00000

Usage

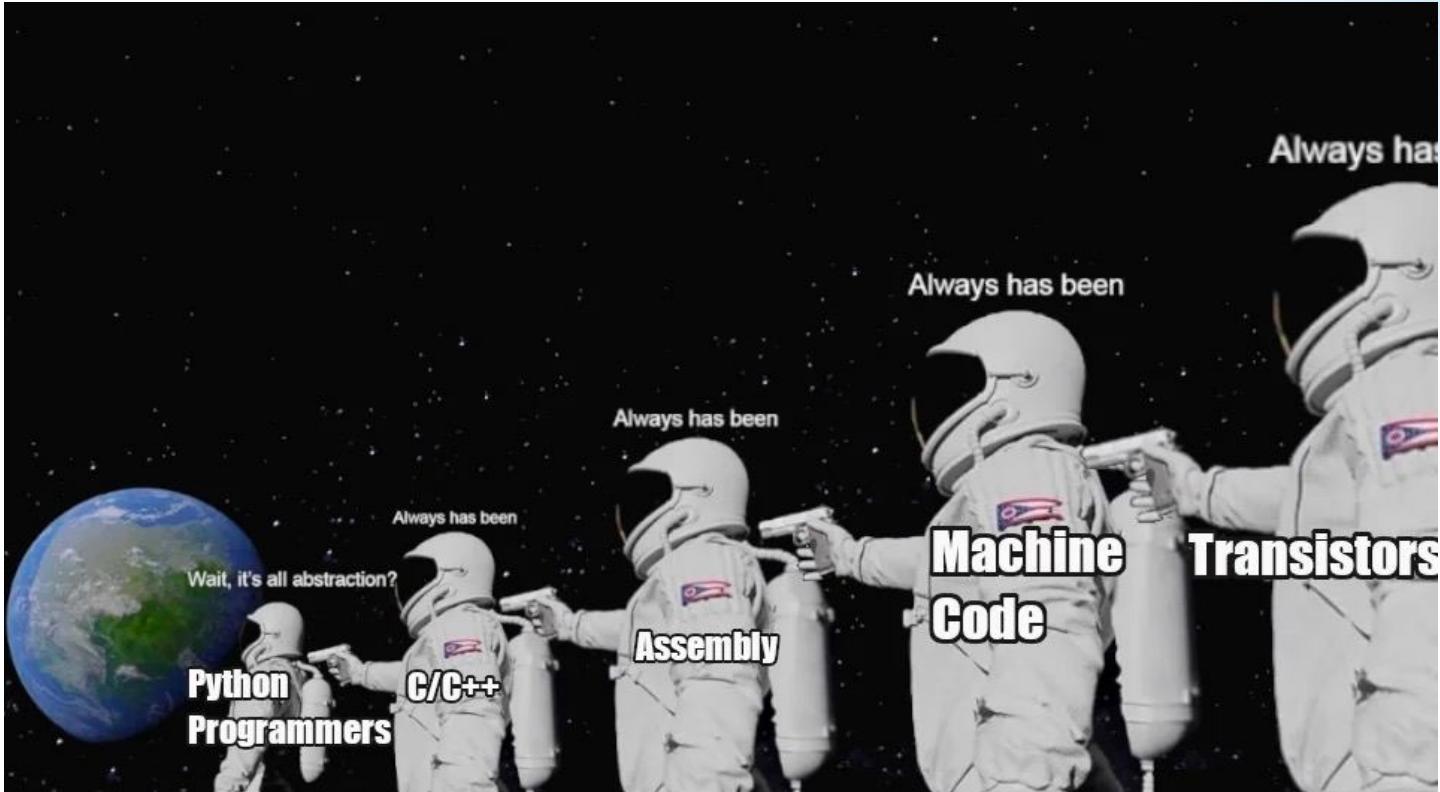
- Numpy
<https://numpy.org/devdocs/reference/simd/index.html>
- Gaming
<https://www.techpowerup.com/forums/threads/avx-vs-non-avx.265592/>
- Many life situations

04

Leaky abstraction

In this section, we will discuss vectorized computations.

Leaky abstraction





EPIC

Institute of Technology
Powered by epam

Vector

The storage of the vector is handled automatically, being expanded as needed. Vectors usually occupy more space than static arrays, because more memory is allocated to handle future growth. This way a vector does not need to reallocate each time an element is inserted, but only when the additional memory is exhausted. The total amount of allocated memory can be queried using [capacity\(\)](#) function. Extra memory can be returned to the system via a call to [shrink_to_fit\(\)](#). (since C++11)

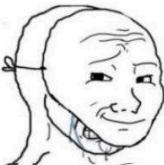


EPIC

Institute of Technology
Powered by epam

TCP vs UDP

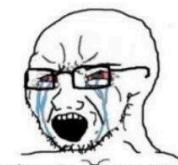
*UDP



Here is your data:

01100100 01010001 01110111 00110100 01110111 00111001 01010111 01100111 01011000 01100011 01010001

No, pls wait... let me..just..



*TCP



would you like to see
a TCP/IP meme?



Yes i would like to see
a TCP/IP meme



Ok, here is the meme
did you receive it?



Yes, i received the meme



Excellent, TCP/IP meme is over.



nice

Premature optimization

“Premature optimization is the root of all evil”

- Knuth

PyPy

Source:

```
1 n = 128
2 d = [[0] * n] * n
3
4 for i in range(n):
5     for j in range(n):
6         d[i][j] = i + j
7
8 for k in range(n):
9     for i in range(n):
10        for j in range(n):
11            d[i][j] = min(d[i][j], d[i][k] + d[k][j])
12
```

Switch off editor

Tab size:

Run

Language:

Input:

Choose file

No more than 256 KB

Output:

```
=====
Used: 46 ms, 2104 KB
```

First 255 bytes only

Without AVX

```
1. #pragma GCC target("avx2")
2. #pragma GCC optimize("O2,unroll-loops")
3.
4. #include <bits/stdc++.h>
5. #include <immintrin.h>
6.
7. using namespace std;
8. const int SIZE = 50000000;
9.
10. double a[SIZE], b[SIZE], c[SIZE];
11.
12. int main() {
13.     for (int i = 0; i < SIZE; i++) {
14.         a[i] = b[i] = i;
15.     }
16.     for (int i = 0; i < SIZE; i++) {
17.         c[i] = a[i] + b[i];
18.     }
19.
20.     return 0;
21. }
```

<https://ideone.com/ceA9iU>

Success #stdin #stdout 0.26s 1175352KB



EPIC

Institute of Technology
Powered by epam

That's All Folks!