

Scheduling

Course: Systems Architecture

Lecturer: Gleb Lobanov

May, 2024



EPIC

Institute of Technology
Powered by epam

Contents

01 Process

02 Protection

03 Threads

04 Race condition

05 Scheduler

06 Fibers

07 GIL



EPIC

Institute of Technology
Powered by epam

01

Processes



EPIC

Institute of Technology
Powered by epam

Processes

USER	PID	%CPU	%MEM	VSZ	RSS	TT	STAT	STARTED	TIME	COMMAND
glebodin	673	11,2	1,3	410311408	223424	??	S	4т06	4:48.31	/Applicatio
_windowserver	362	7,4	0,7	411053248	123856	??	Ss	4т06	160:06.67	/System/Lib
glebodin	739	4,6	0,4	442387360	71472	??	S	4т06	8:43.04	/Applicatio
glebodin	4935	2,3	1,3	1598023296	226272	??	S	11:30	36:54.08	/Applica
glebodin	1904	2,0	0,0	408304000	7600	s000	S	7:27	0:10.48	-zsh
glebodin	663	1,6	2,8	443056784	471312	??	S	4т06	21:50.40	/Applicatio
glebodin	16795	1,5	0,3	408599296	46256	??	S	7:05	0:00.15	/System/L
glebodin	2888	1,3	1,9	10423780	310688	??	S	8:40	37:14.66	/Users/gle
glebodin	850	1,0	0,6	10731028	93072	??	S	4т06	5:51.41	/opt/homebre
glebodin	2899	0,9	0,9	10209980	157136	??	Ss	8:40	11:42.95	/Users/gle
glebodin	12056	0,9	1,7	410224016	289344	??	S	2:02	57:40.31	/Applicat
root	323	0,5	0,1	409186000	10592	??	Ss	4т06	1:43.82	/System/Lib
glebodin	16673	0,5	0,1	409183104	23008	??	S	7:01	0:00.55	/System/L
root	449	0,5	0,1	408996432	22224	??	Ss	4т06	5:43.93	/usr/libexe
glebodin	670	0,5	2,7	417848288	458752	??	S	4т06	66:15.91	/Applicatio
glebodin	16794	0,4	0,2	408606272	37696	??	S	7:05	0:00.14	/System/L
root	308	0,3	0,2	409905424	28912	??	Ss	4т06	4:39.81	/System/Lib
glebodin	2922	0,2	2,5	14386044	426816	??	S	8:40	2:17.90	/Users/gle
_networkd	436	0,2	0,1	408635696	14400	??	Ss	4т06	0:18.64	/usr/libexe
root	532	0,2	0,5	414207504	86160	??	Ss	4т06	3:35.34	/System/Lib
glebodin	16752	0,2	0,0	408067040	4720	??	S	7:04	0:00.22	/System/L
root	281	0,1	0,0	408632224	6752	??	Ss	4т06	1:31.28	/System/Lib
root	510	0,1	0,3	409251792	51808	??	SN	4т06	22:11.07	/opt/osquer
root	292	0,1	0,2	408941920	25840	??	Ss	4т06	5:38.02	/Library/AP
glebodin	16698	0,1	0,1	409001184	19632	??	Ss	7:02	0:00.19	/System/L
glebodin	4926	0,1	0,5	1585149744	89792	??	S	11:30	1:04.29	/Applica
glebodin	5543	0,1	1,1	441033632	180832	??	S	11:53	9:37.72	/Applicat
_locationd	16729	0,0	0,0	408361504	5152	??	S	7:04	0:00.06	/usr/libe
root	16728	0,0	0,0	408361488	5024	??	Ss	7:04	0:00.05	/usr/libe
glebodin	16727	0,0	0,1	408630528	11296	??	S	7:04	0:00.15	/usr/libe
glebodin	16726	0,0	0,1	408635648	14422	??	S	7:04	0:00.25	/System/L



EPIC

Institute of Technology
Powered by epam

Let's touch them

```
import itertools
import psutil

# getting all processes
all_processes = psutil.process_iter()

# getting 5 processes
limited_processes = itertools.islice(all_processes, 5)

# output info about processes
for process in limited_processes:
    pid = process.pid
    name = process.name()
    print(f"PID: {pid}, Name: {name}")

process = psutil.Process(16141)
print("process id:", process.pid)
print("process name:", process.name())
print("cp usage:", process.cpu_percent())
print("usage of memory:", process.memory_info().rss)
print("status:", process.status())
print("time of creation:", process.create_time())
print("parent id:", process.ppid())
```

Let's touch them

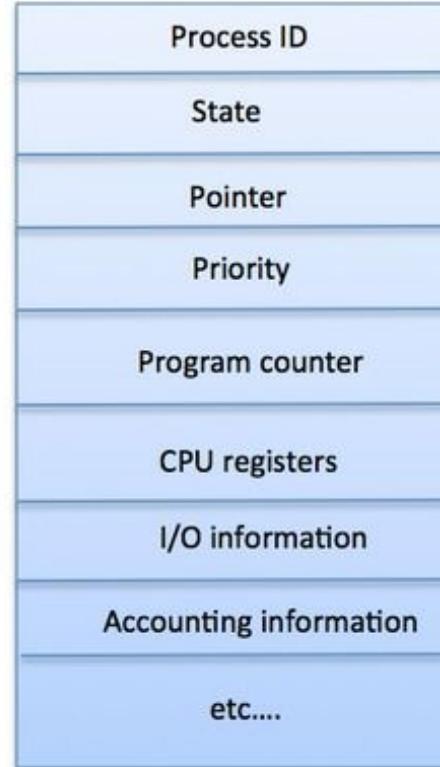
```
└─ python3 processes.py
    PID: 0, Name: kernel_task
    PID: 1, Name: launchd
    PID: 277, Name: syslogd
    PID: 278, Name: UserEventAgent
    PID: 280, Name: uninstalld
    process id: 16141
    process name: deleted
    cp usage: 0.0
    usage of memory: 8601600
    status: running
    time of creation: 1685115830.214878
    parent id: 1
```



EPIC

Institute of Technology
Powered by epam

Process info



Priority - NI. The smaller the "nice" value, the higher the priority of the process.

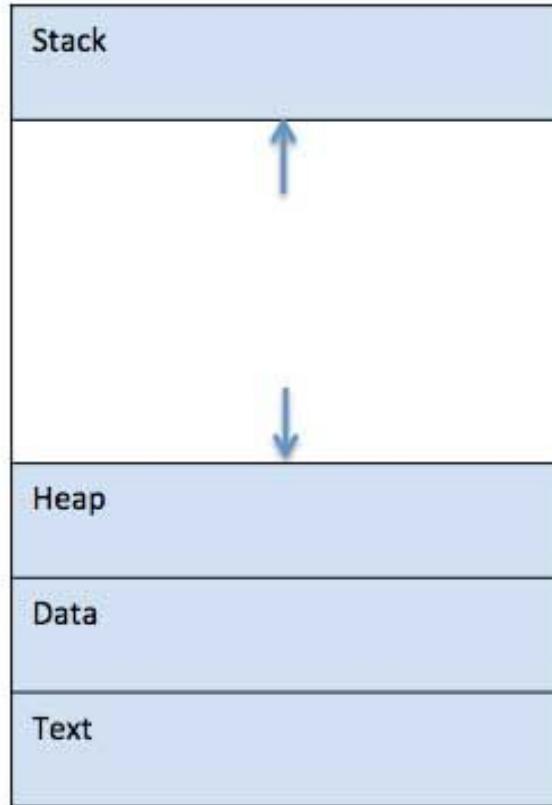
F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
24	1801	S	200	7032	7286	0	60	20	1b4c	108		pts/2	0:00 ksh
200	801	S	200	7568	7032	0	70	25	2310	88	5910a58	pts/2	0:00 vmstat
24	1801	S	200	8544	6494	0	60	20	154b	108		pts/0	0:00 ksh



EPIC

Institute of Technology
Powered by epam

Process stack

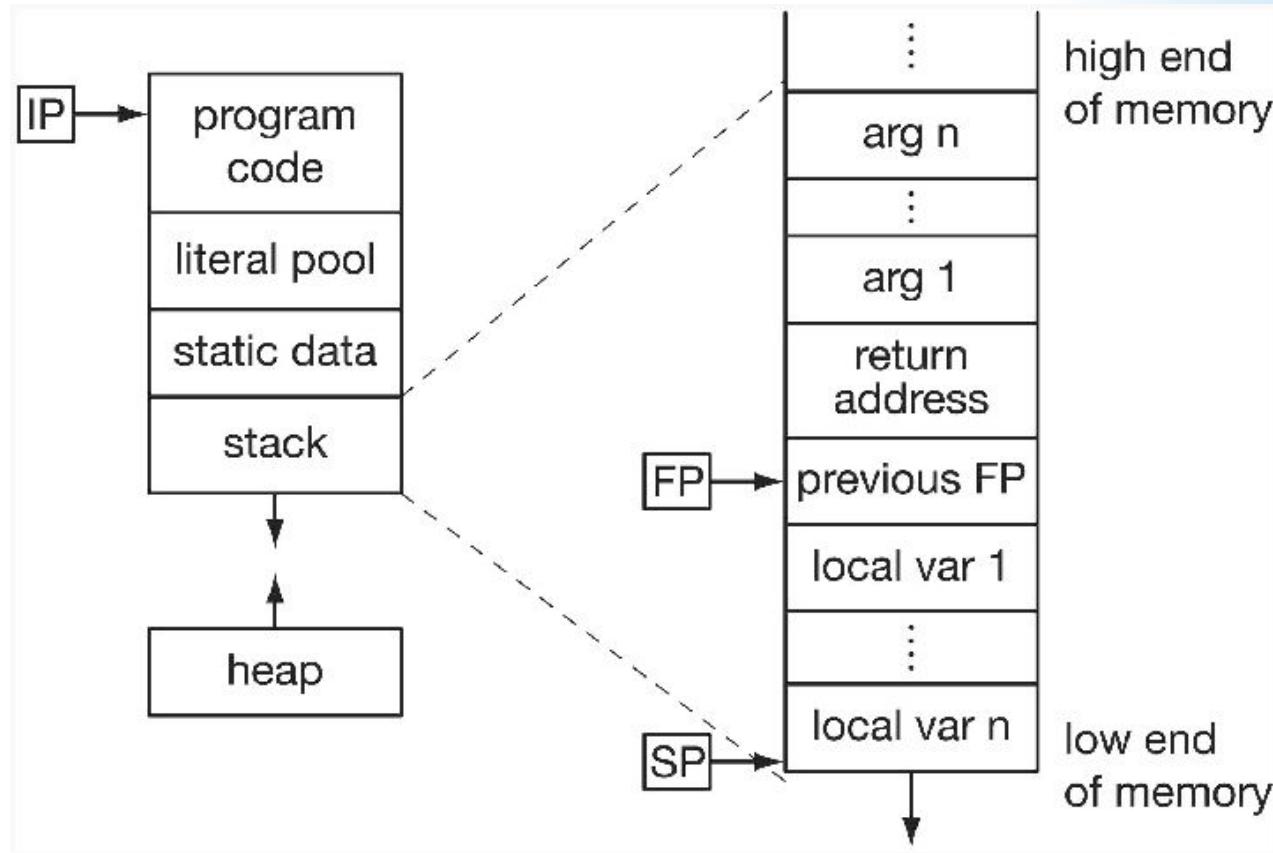




EPIC

Institute of Technology
Powered by epam

Context



Component context

- 1) General Purpose Registers
- 2) Command Counter
- 3) Stack context

Process

Process Control Block (PCB)

Facilities for process

Pointer Process state	
Process ID	Unique ID
Program counter	Next program that run
Registers	
Memory Limits	
Accounting	Log INFO about process
List of open file	

PCB Diagram

C++

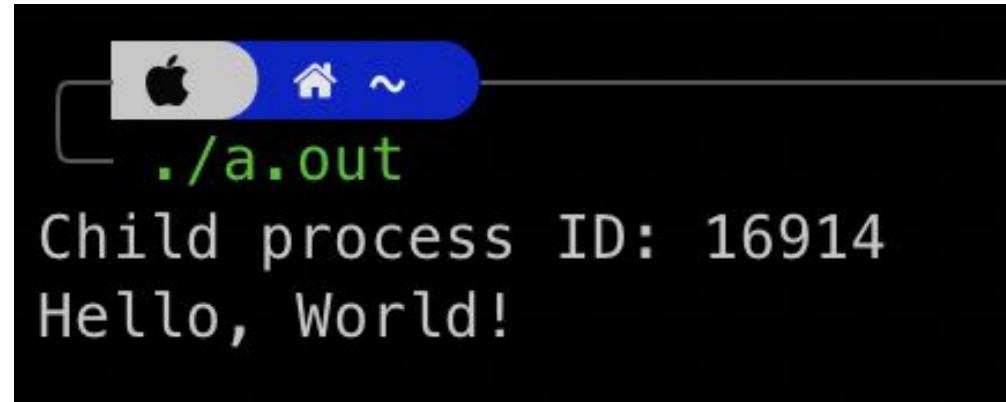
not perfect)

```
#include <iostream>
#include <cstdlib>
#include <unistd.h>

int main() {
    // Creating a new process
    pid_t pid = fork();

    if (pid == -1) {
        // Handling error in process creation
        std::cerr << "Error creating a process!" << std::endl;
        return 1;
    } else if (pid == 0) {
        // Code executed in the child process
        // Printing "Hello, World!"
        std::cout << "Hello, World!" << std::endl;
    } else {
        // Code executed in the parent process
        // Outputting child process ID
        std::cout << "Child process ID: " << pid << std::endl;
        // Waiting for the child process to finish
        int status;
        waitpid(pid, &status, 0);
    }
    return 0;
}
```

Result

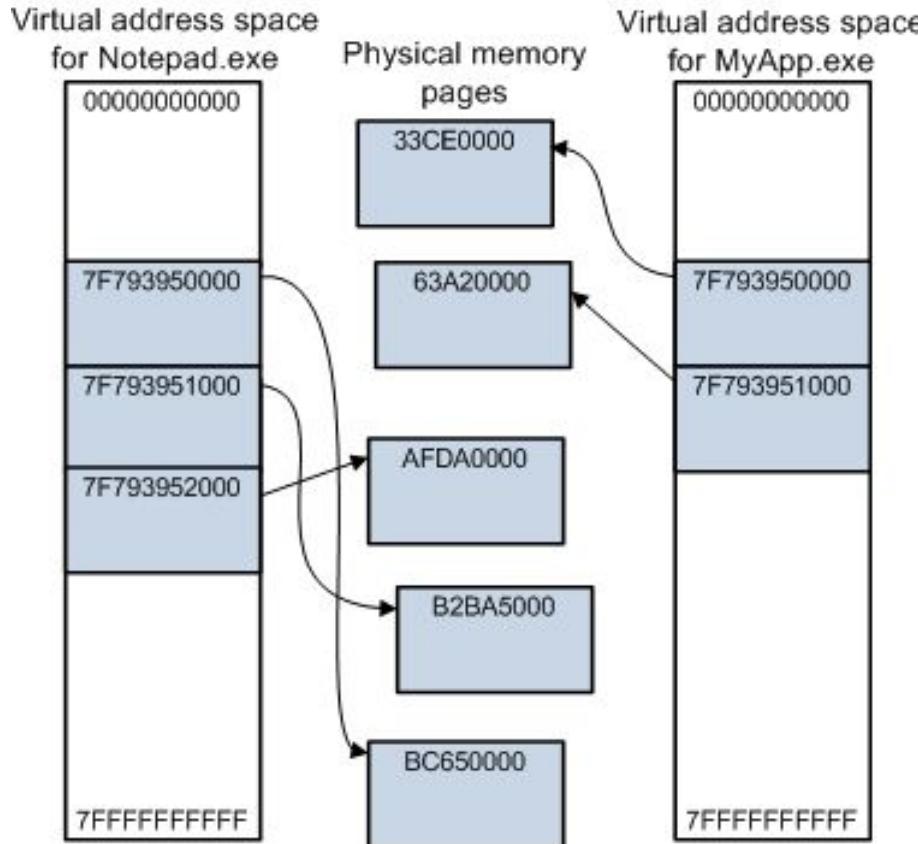


```
./a.out
Child process ID: 16914
Hello, World!
```

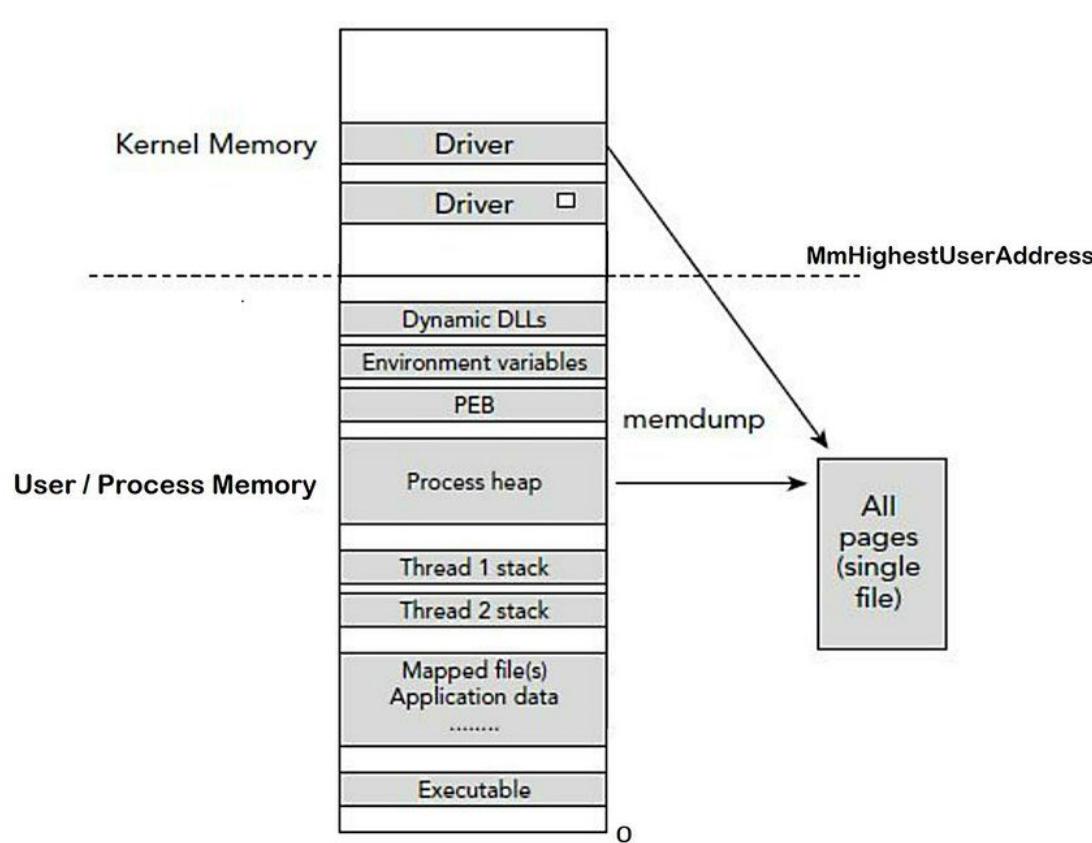
02

Protection

Virtual address space



Virtual address space



Protection

MPROTECT(2)

BSD System Calls Manual

MPROTECT(2)

NAME

mprotect -- control the protection of pages

SYNOPSIS

```
#include <sys/mman.h>
```

```
int  
mprotect(void *addr, size_t len, int prot);
```

DESCRIPTION

The **mprotect()** system call changes the specified pages to have protection *prot*. Not all implementations will guarantee protection on a page basis but Mac OS X's current implementation does.

When a program violates the protections of a page, it gets a SIGBUS or SIGSEGV signal.

Currently *prot* can be one or more of the following:

PROT_NONE	No permissions at all.
PROT_READ	The pages can be read.
PROT_WRITE	The pages can be written.
PROT_EXEC	The pages can be executed.

RETURN VALUES

Upon successful completion, a value of 0 is returned. Otherwise, a value of -1 is returned and *errno* is set to indicate the error.

ERRORS

mprotect() will fail if:

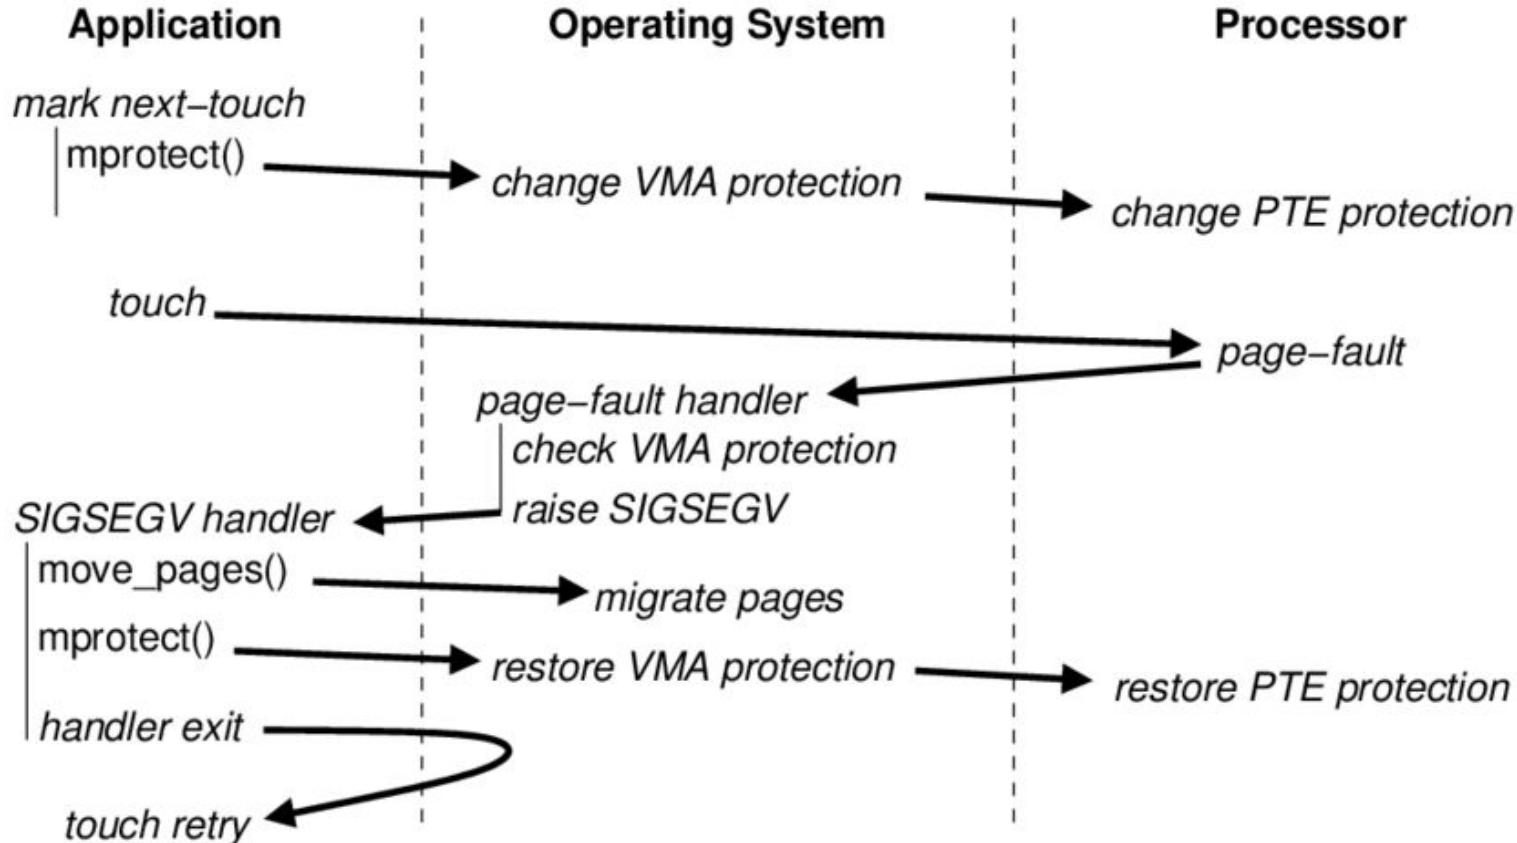
[EACCES]	The requested protection conflicts with the access permissions of the process on the specified address range.
[EINVAL]	<i>addr</i> is not a multiple of the page size (i.e. <i>addr</i> is not page-aligned).
[ENOMEM]	The specified address range is outside of the address range of the process or includes an unmapped page.
[ENOTSUP]	The combination of accesses requested in <i>prot</i> is not supported.



EPIC

Institute of Technology
Powered by epam

Protection



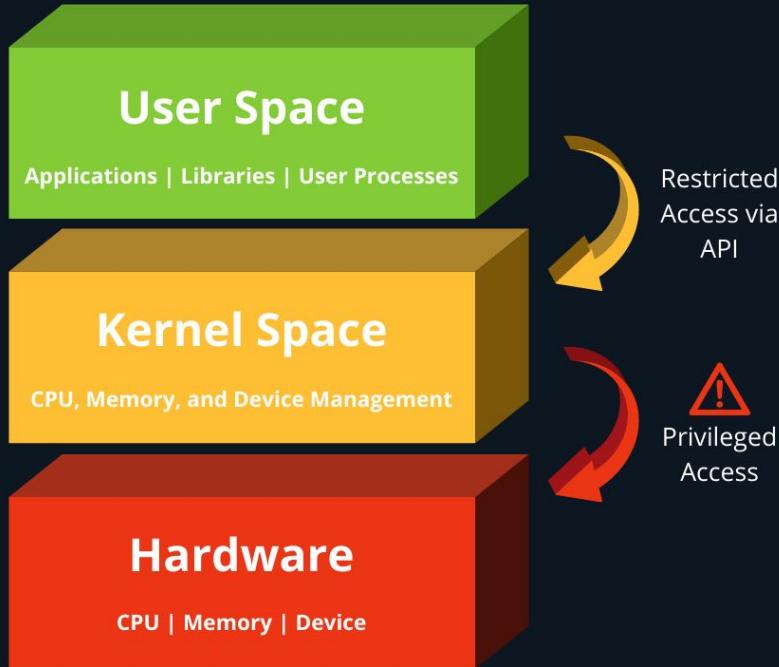


EPIC

Institute of Technology
Powered by epam

Modes

Differences Between Kernel and User Mode



User Mode

User mode agents run in the same space as common user applications. Lakeside's agent operates in user mode gathering thousands of metrics in real-time on performance and user behavior.

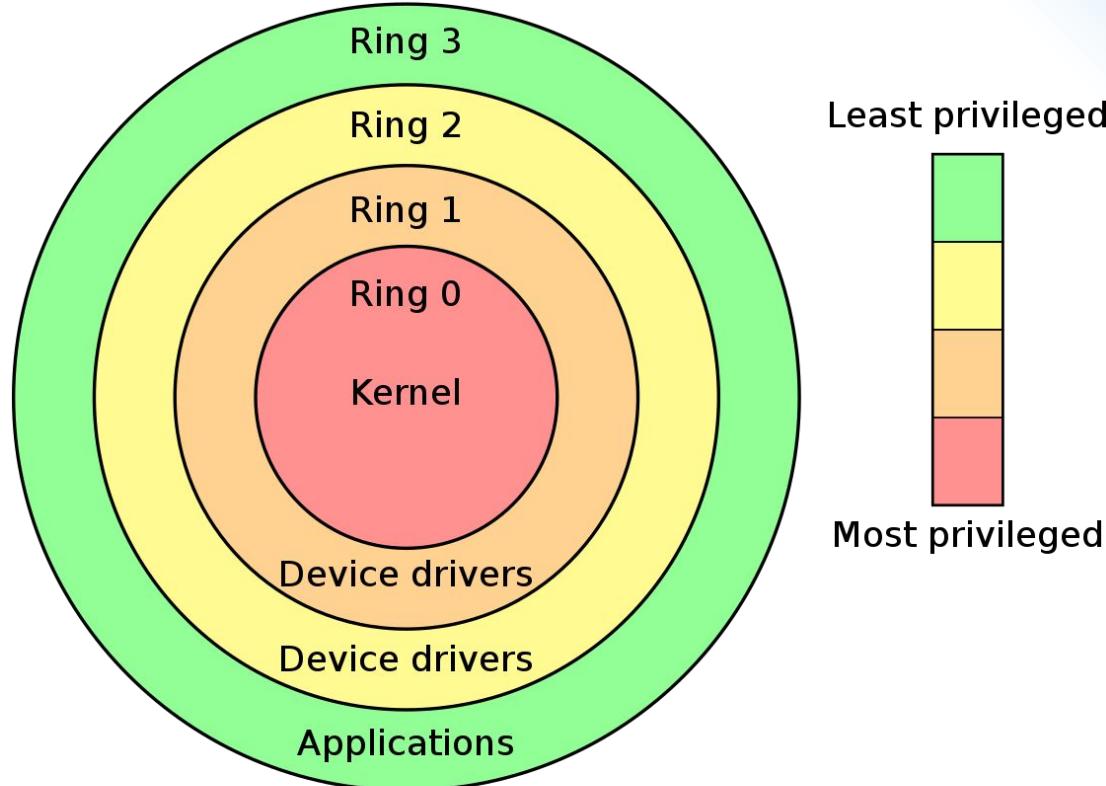
Kernel Mode

Kernel mode agents operate beneath the operating system with privileged access to bare hardware. Risks associated with kernel mode agents include cyberattacks and system crashes.



Lakeside®

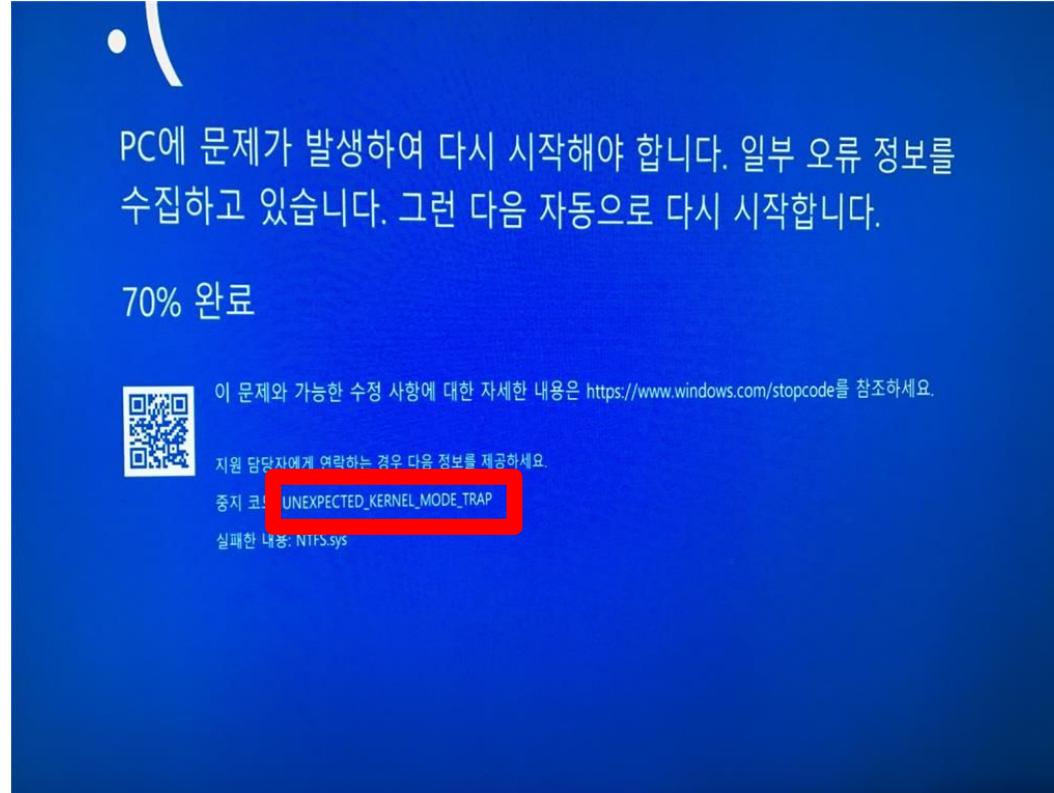
Rings



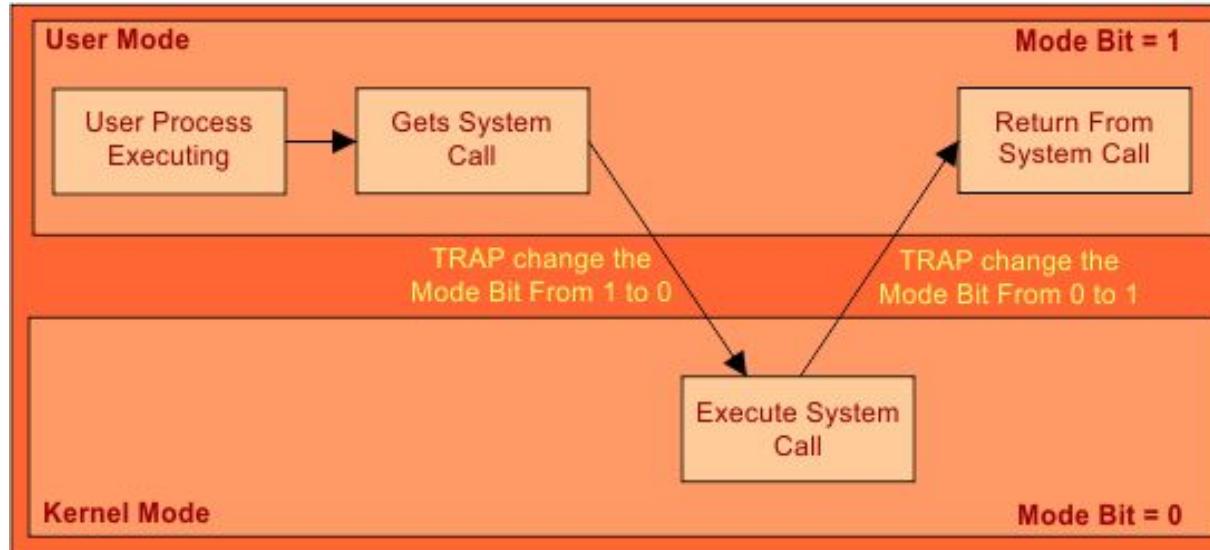
Access of processes

- 1) Separation of execution modes
- 2) Permissions to access memory pages
- 3) Syscalls
- 4) Security modules and policies

Blue screen

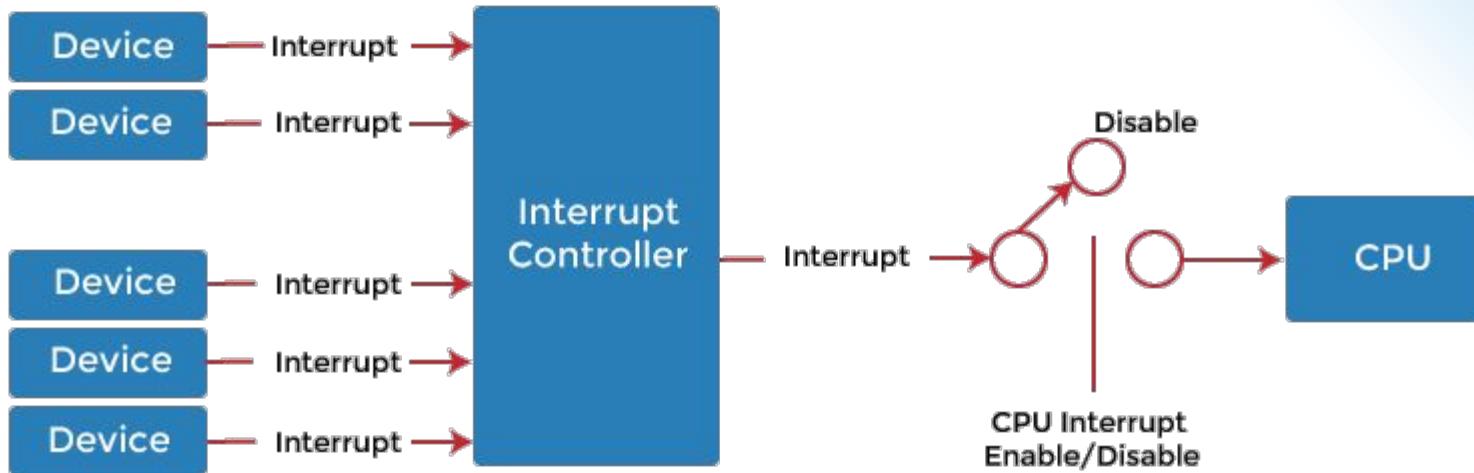


Syscalls

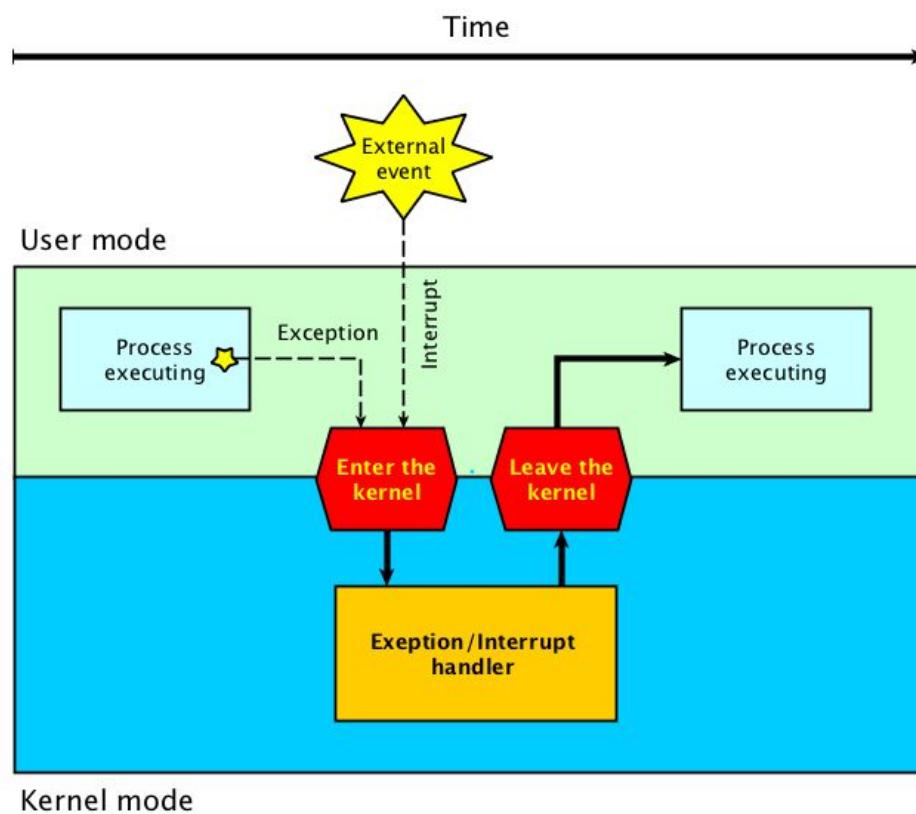




Interrupt



Interrupt handling



Interrupt types

- 1) Hardware Interrupts
- 2) Software Interrupts
- 3) External Interrupts
- 4) Timer Interrupts
- 5) Virtual Interrupts

Timer interrupts types

- 1) Task Scheduling
- 2) System time update
- 3) Operations with precise time



EPIC

Institute of Technology
Powered by epam

Interrupt

```
#include <iostream>

int main() {
    std::cout << "Before interrupt." << std::endl;

    // Calling a software interrupt
    asm volatile("int $3"); // Interrupt number 3

    std::cout << "After interrupt." << std::endl;

    return 0;
}
```



EPIC

Institute of Technology
Powered by epam

03

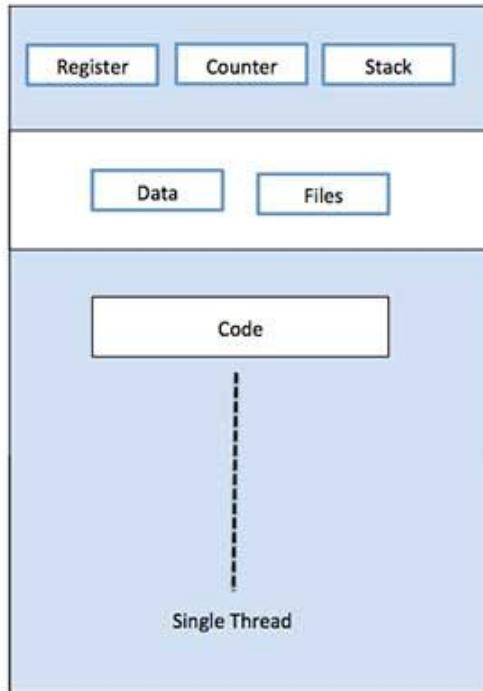
Threads



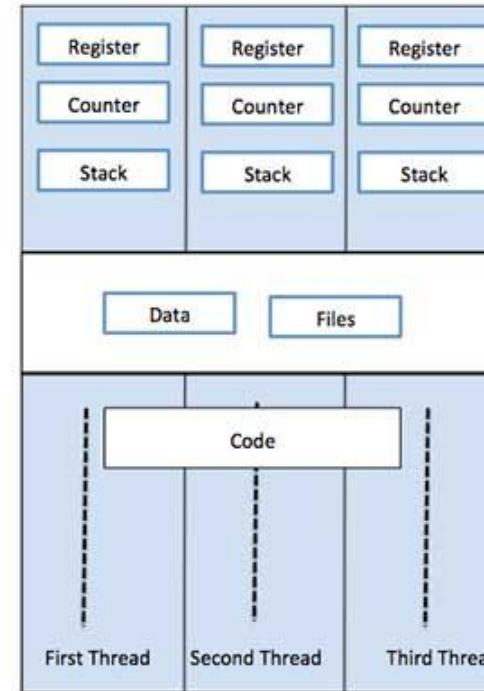
EPIC

Institute of Technology
Powered by epam

Threads



Single Process P with single thread



Single Process P with three threads

Process vs Thread

Classification	Process	Threads
Definition	Execution of any program is a process	Segment or subset of a process is a thread
Communication	It takes more time to communicate between the different processes.	It takes less time for the communications
Resources	It consumes more resources	It consumes fewer resources
Memory	The process does not share the memory and is carried out alone	Threads are used to share their memory
Data Sharing	Process does not share their data	Threads share their data
Size	The process consumes more size	Threads are lightweight
Execution error	One process will not affect any other process due to an error	Threads will not run if one thread has an error
Termination time	Process usually takes more time to terminate	Threads take less time to terminate

Threads code

```
#include <iostream>
#include <thread>

// Function to be executed by the thread
void threadFunction() {
    std::cout << "This is a thread." << std::endl;
}

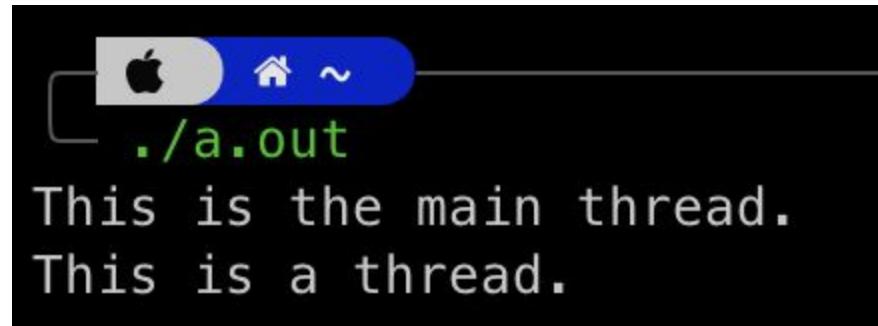
int main() {
    // Creating a thread object and passing it the function to execute
    std::thread myThread(threadFunction);

    std::cout << "This is the main thread." << std::endl;

    // Waiting for the thread to finish
    myThread.join();

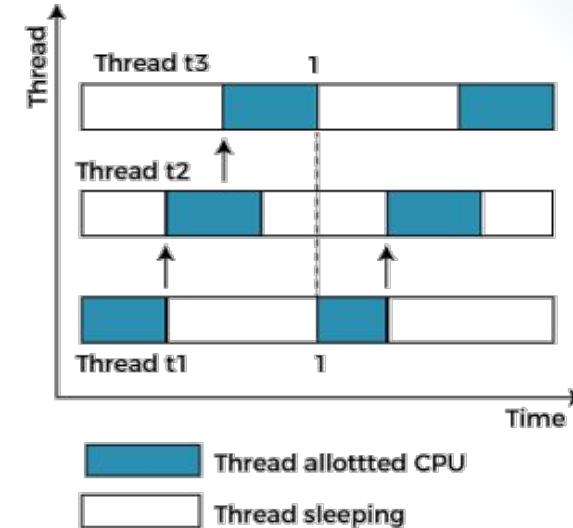
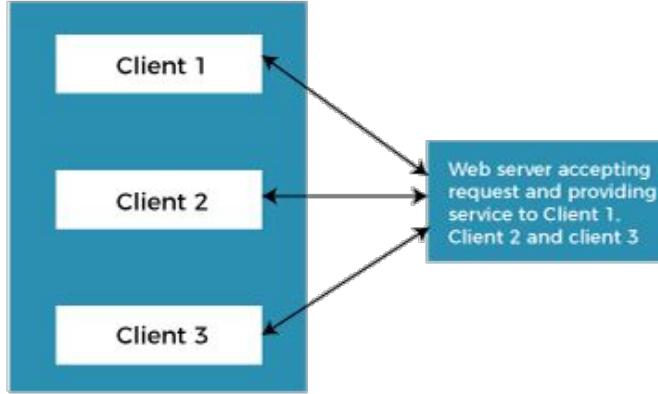
    return 0;
}
```

Result



A screenshot of a macOS terminal window. The window title bar shows the Apple logo, a blue rounded rectangle with a white house icon and a tilde (~), and a close button. The terminal prompt is ./a.out. The output text is:
This is the main thread.
This is a thread.

Threads control



Threads control in OS

- 1) Creating streams
- 2) Scheduling flows
- 3) Context switching
- 4) Synchronization and interaction
- 5) Priority Management
- 6) Ending Flows

Why we can't make it simple?

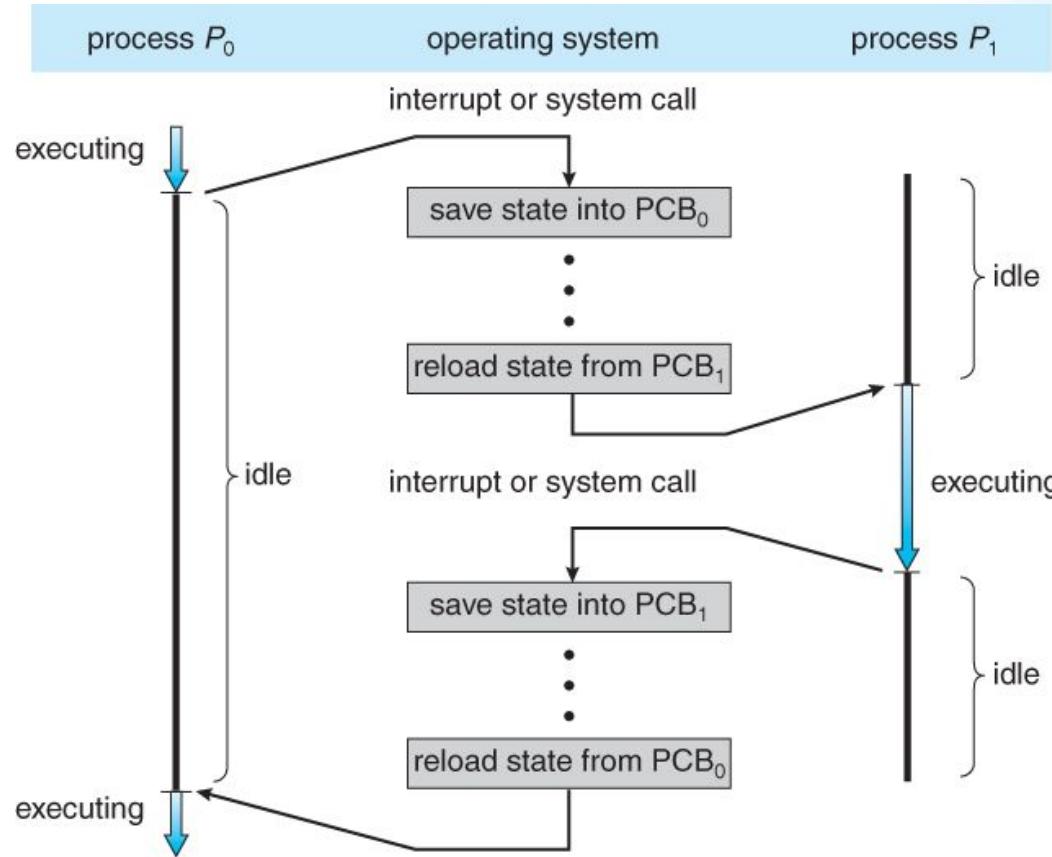
- 1) Poor use of resources
- 2) Parallel computing
- 3) Responsiveness
- 4) Synchronization and collaboration
- 5) Multitasking support



EPIC

Institute of Technology
Powered by

Switching



How to change task?

- 1) Cooperative multitasking
- 2) Interruptions by time interval
- 3) Event-based interruptions
- 4) I/O multiplexing

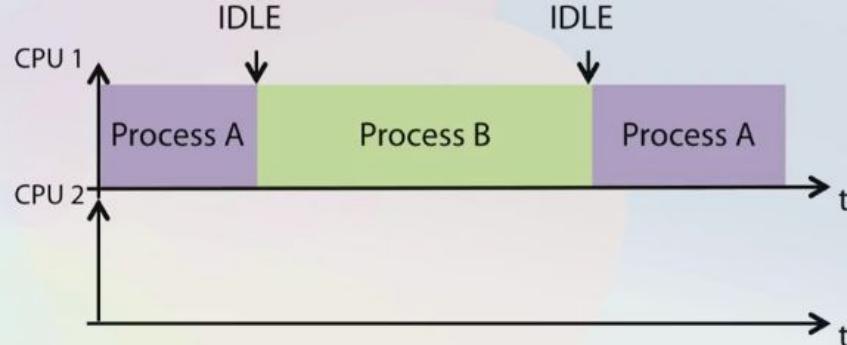


EPIC

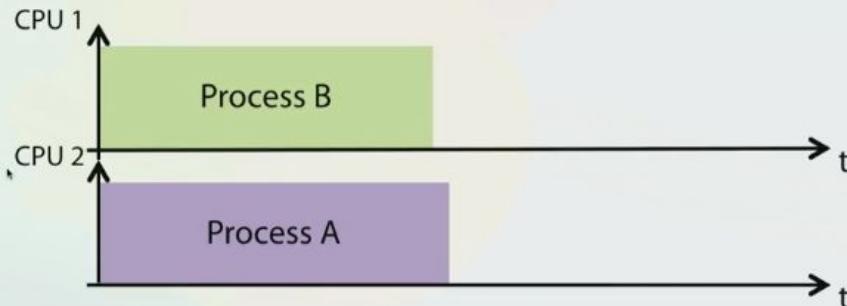
Institute of Technology
Powered by epam

Multitasking

Cooperative mode

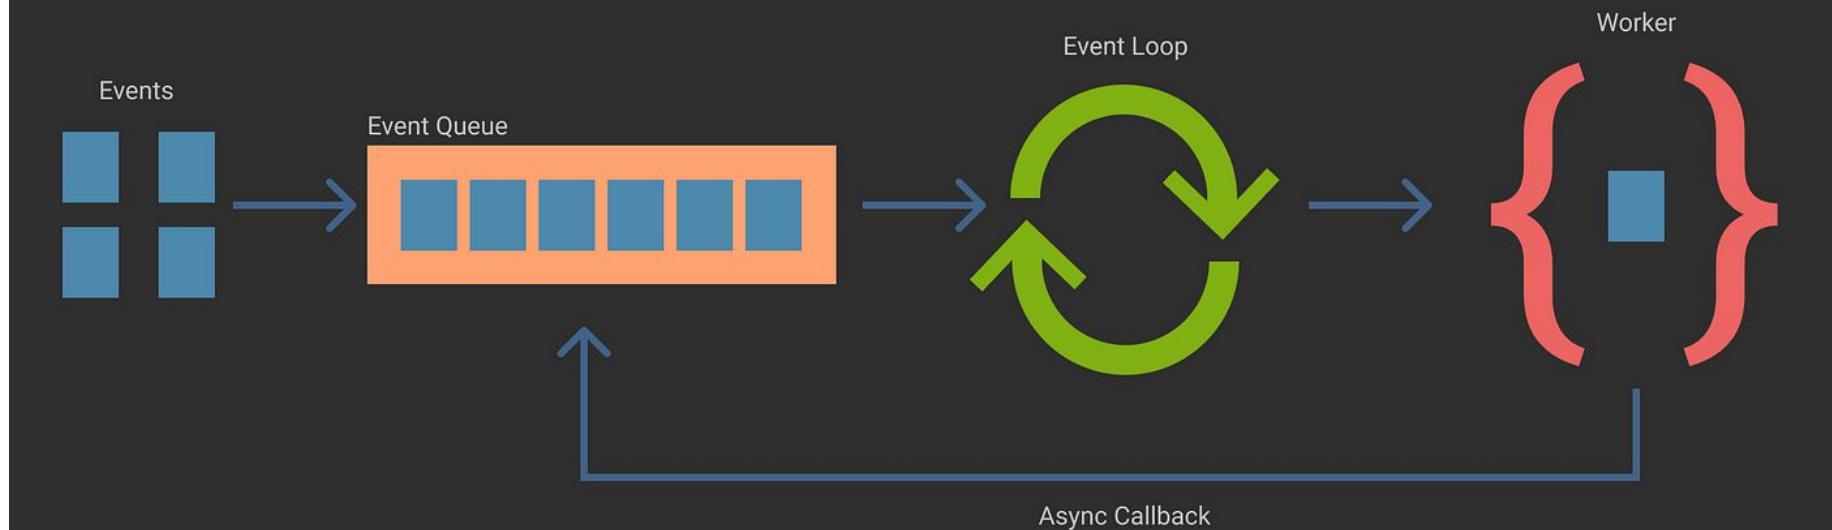


Preemptive mode



Event loop

Event Loop With Single Worker





EPIC

Institute of Technology
Powered by epam

04

Race condition



Race condition

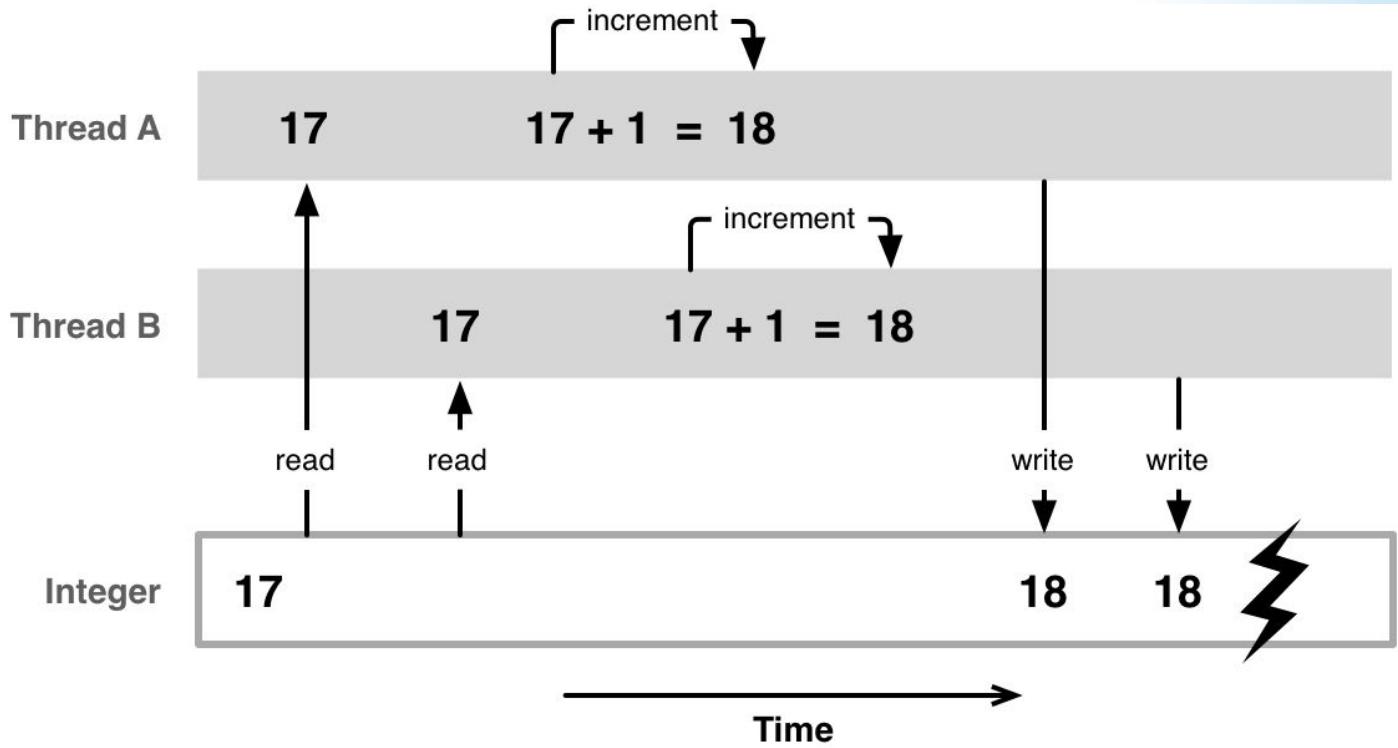




EPIC

Institute of Technology
Powered by

Race condition



Bug bounty

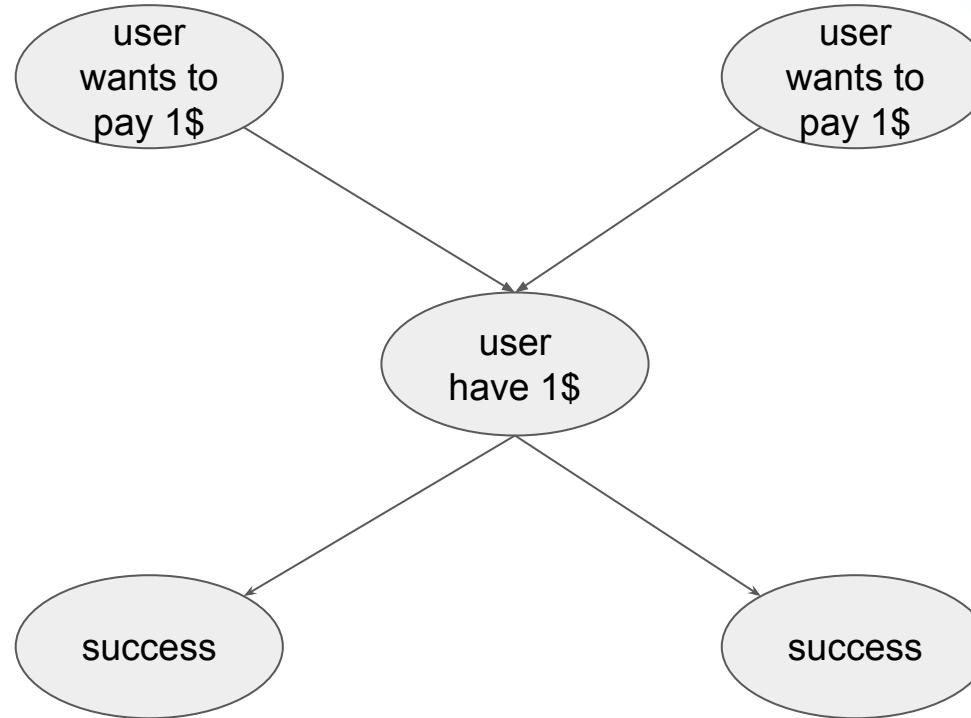
Most of OAuth 2 API implementations seem to have multiple Race Condition vulnerabilities for processing requests for Access Token or Refresh Token.

Race Condition allows a malicious application to obtain several `access_token` and `refresh_token` pairs while only one pair should be generated. Further, it leads to authorization bypass when access would be revoked.

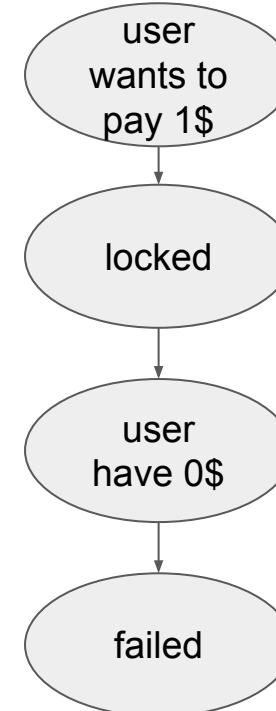
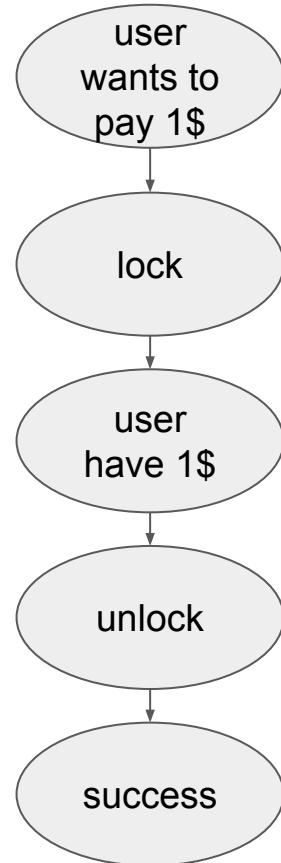
I've already tested for this vulnerability 11 different targets (web-services providing OAuth2 API), and 6 of them are vulnerable. Only one target seems to be certainly protected (or I just failed to catch into Race Condition time window). The rest 4 targets have Race Condition bug, but protected against further exploitation by one of the following reasons:

- only one of several tokens generated is valid (**1 target**)
- for any access revocation all tokens always are revoked (**1 target**)
- for all concurring requests finished successfully the same `access_token` values (or in pair with `refresh_token`) obtained (**2 targets**)

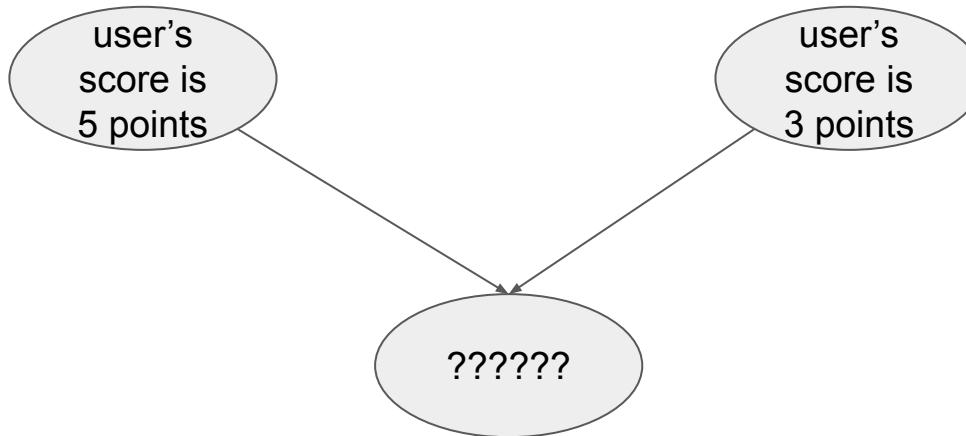
Banking



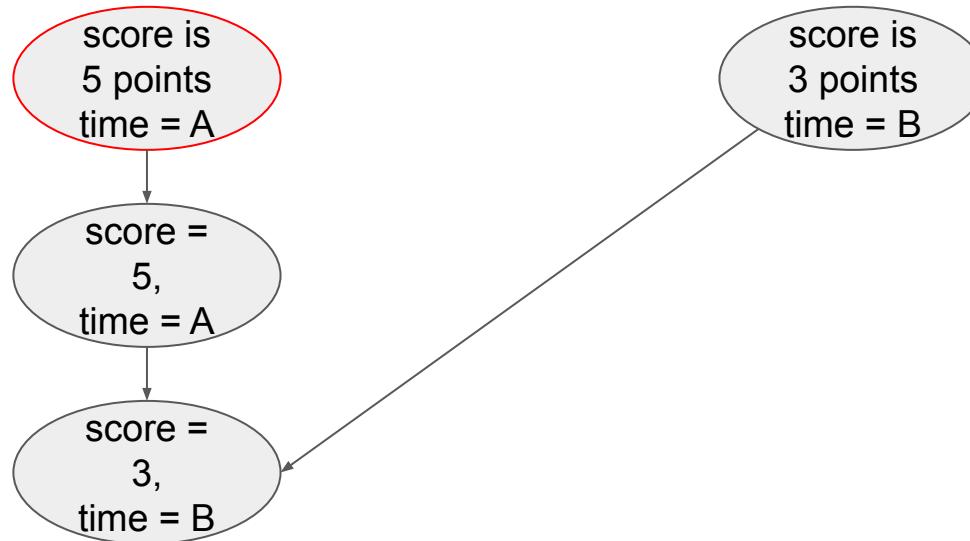
Solution



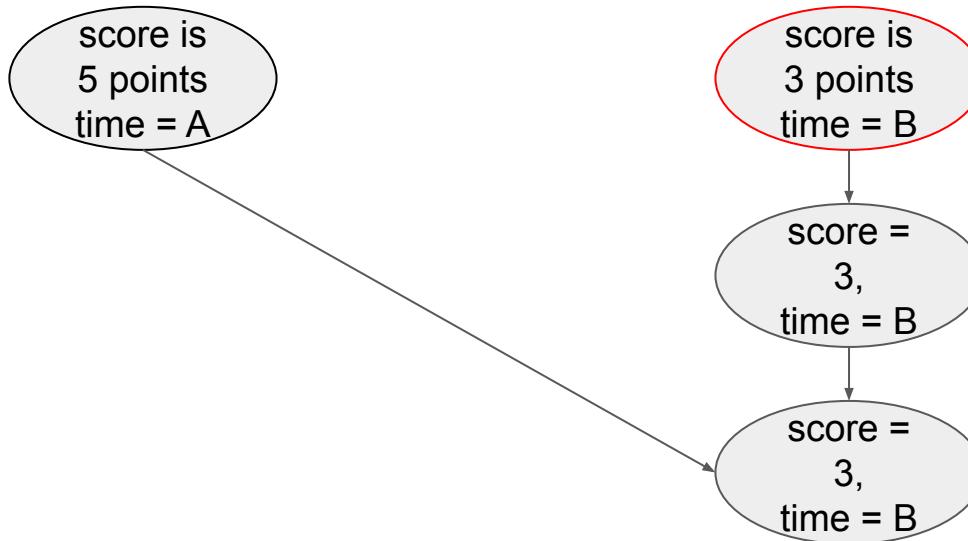
Game



Solution



Solution



Atomicity

Atomic operation : is an indivisible operation. We cannot observe such an operation half-done from any thread in the system.

i++;

- Step 1 : Read the value of 'i' from memory
- Step 2 : Increment that value by one
- Step 3 : Store the modified value to the memory

Thread 1



Thread 2





EPIC

Institute of Technology
Powered by epam

Atomicity

```
#include <atomic>
#include <iostream>
#include <thread>

std::atomic<int> counter; // Atomic variable

void IncrementCounter()
{
    for (int i = 0; i < 1000000; ++i)
    {
        counter.fetch_add(1); // Atomic increment operation
    }
}

int main()
{
    counter = 0;

    std::thread t1(IncrementCounter);
    std::thread t2(IncrementCounter);

    t1.join();
    t2.join();

    std::cout << "Counter value: " << counter << std::endl;

    return 0;
}
```



EPIC

Institute of Technology
Powered by epam

Atomicity

```
└─ [?] ─ Apple ─ ┌ ~/Desktop
   └─ g++ thread_example.cpp -o a.out -std=c++17
```

```
└─ [?] ─ Apple ─ ┌ ~/Desktop
   └─ ./a.out
```

Counter value: 2000000



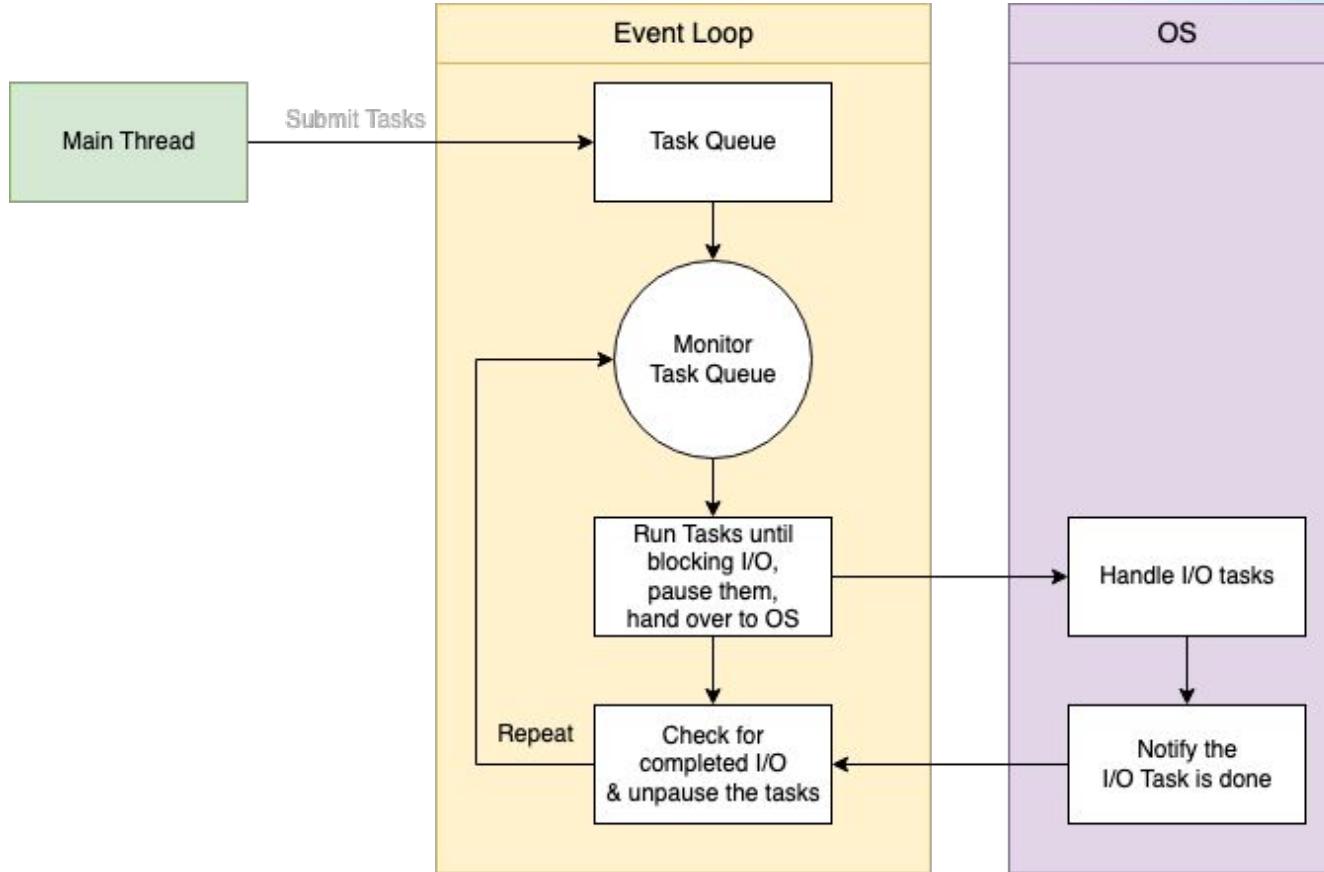
EPIC

Institute of Technology
Powered by epam

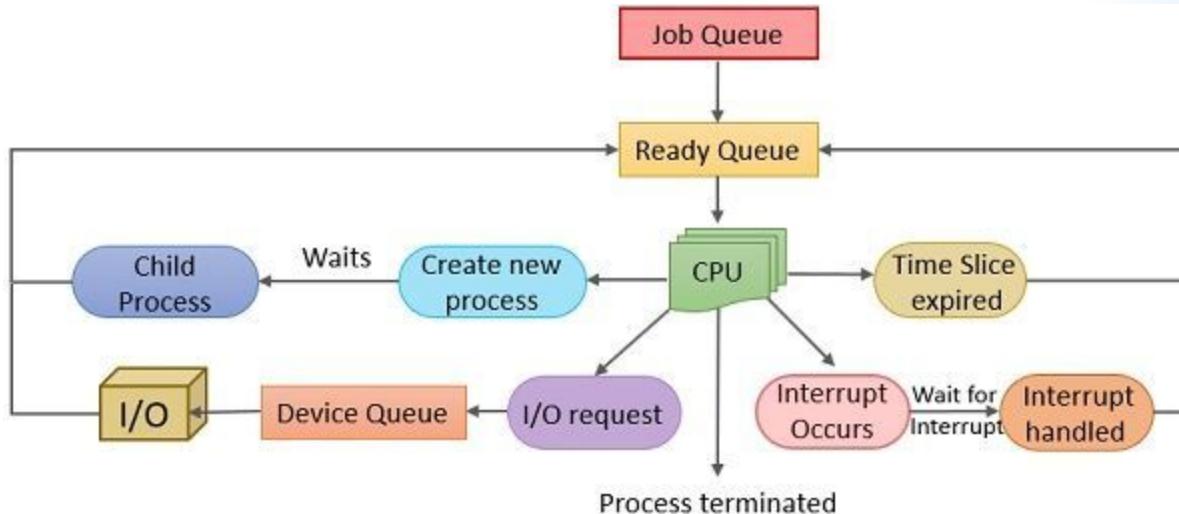
05

Scheduling

Sheduler



Scheduling



Scheduling Queue in Operating System

Realizations

- 1) Round Robin
- 2) Priority Scheduling
- 3) Multilevel Queue
- 4) Earliest Deadline First
- 5) Adaptive Scheduling



EPIC

Institute of Technology
Powered by epam

Priority

Priority Scheduling Algorithm

At time t=16ms

0 1 2 5 10 16



TIME



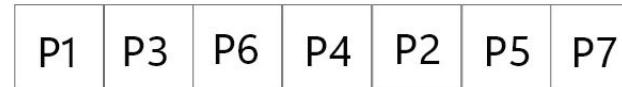
Process	P1	P2	P3	P4	P5	P6	P7
Burst Time	1	7	3	6	5	15	8
Priority	2	6	3	5	4	10	9
Arrival Time	0	1	2	3	4	5	15



Priority

Priority Scheduling Algorithm

0 3 7 11 13 18 27 37



Process	P1	P2	P3	P4	P5	P6	P7
Burst Time	3	5	4	2	9	4	10
Priority	3	6	3	5	7	4	10
Arrival Time	0	2	1	4	6	5	7

Priority Scheduling

PROS

1. Higher priority - faster solution
2. Scales well

CONS

1. Starvation
2. The demands of good priorities

Round-Robin

P1	P2	P3	P4	P5	P6	P1	P3	P4	P5	P6	P3	P4	P5	66
0	5	9	14	19	24	29	31	36	41	46	50	55	56	

Process ID	Arrival Time	Burst Time	Completion Time	Turn Around Time	Waiting Time
P1	0	7	31	31	24
P2	1	4	9	8	4
P3	2	15	55	53	38
P4	3	11	56	53	42
P5	4	20	66	62	42
P6	4	9	50	46	37

Round-Robin

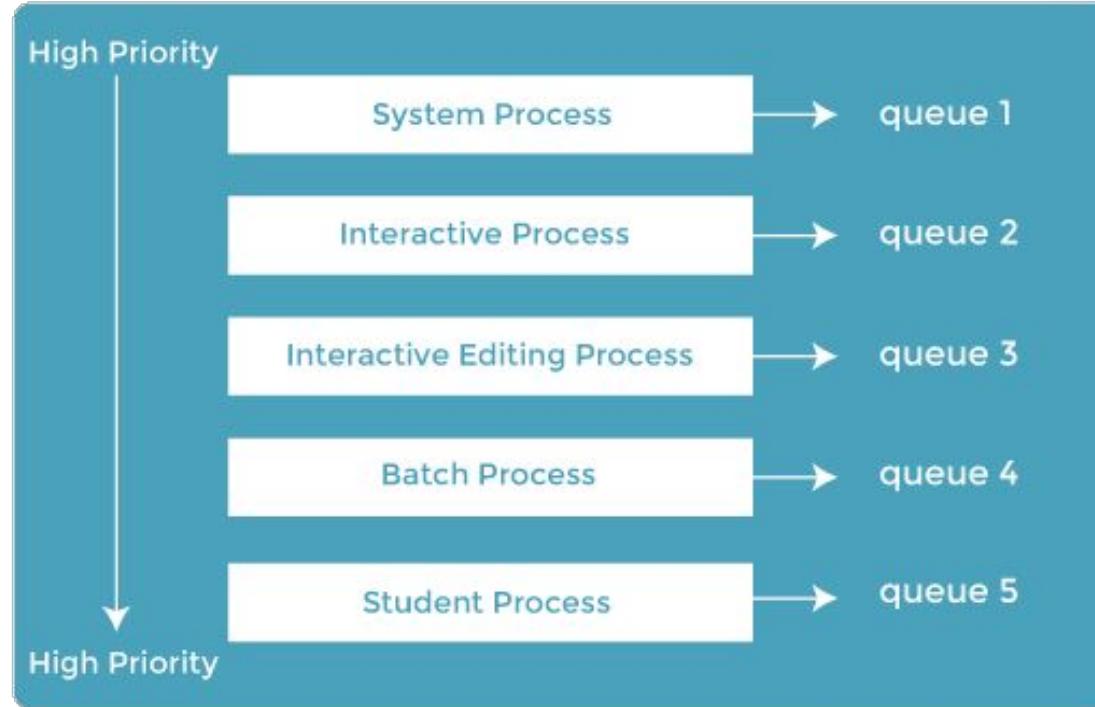
PROS

1. Fair distribution of time
2. Relatively simple implementation
3. Efficient use of CPU resources

CONS

1. Doesn't work well for different work times
2. Additional overhead for context switching between tasks

Multilevel Queue



Multilevel Queue

PROS

1. Balance between different priorities
2. Work differentiation

CONS

1. Starvation
2. A lot of problems with the setup



EPIC

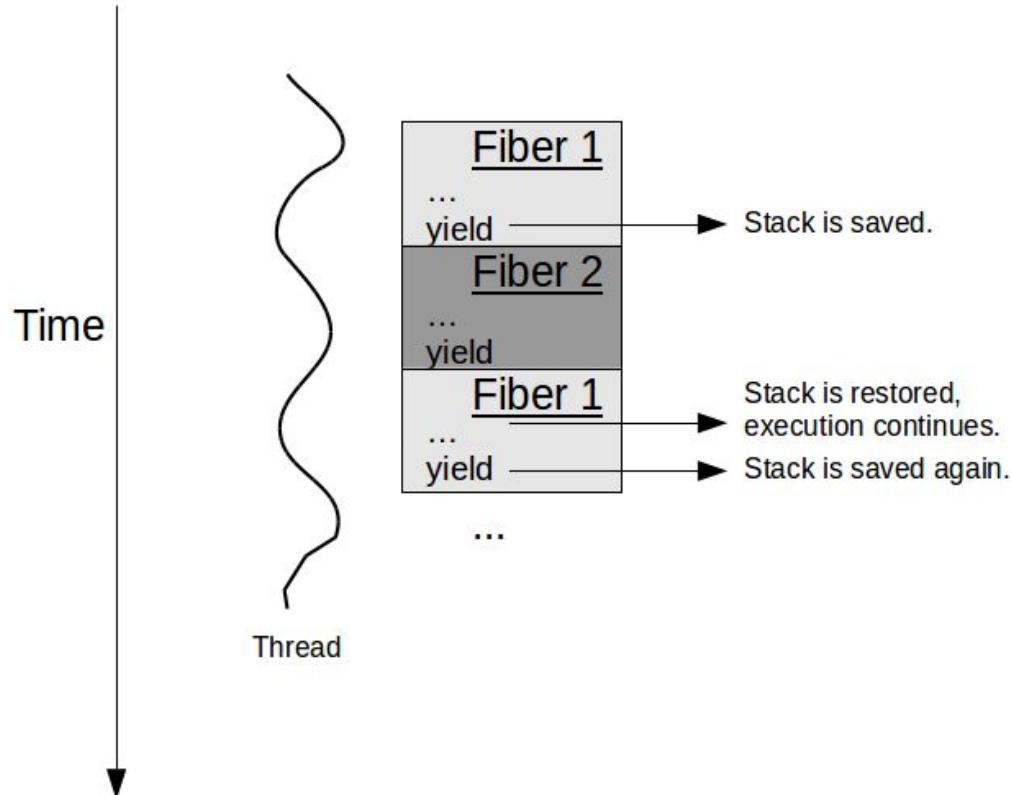
Institute of Technology
Powered by epam

06

Fiber



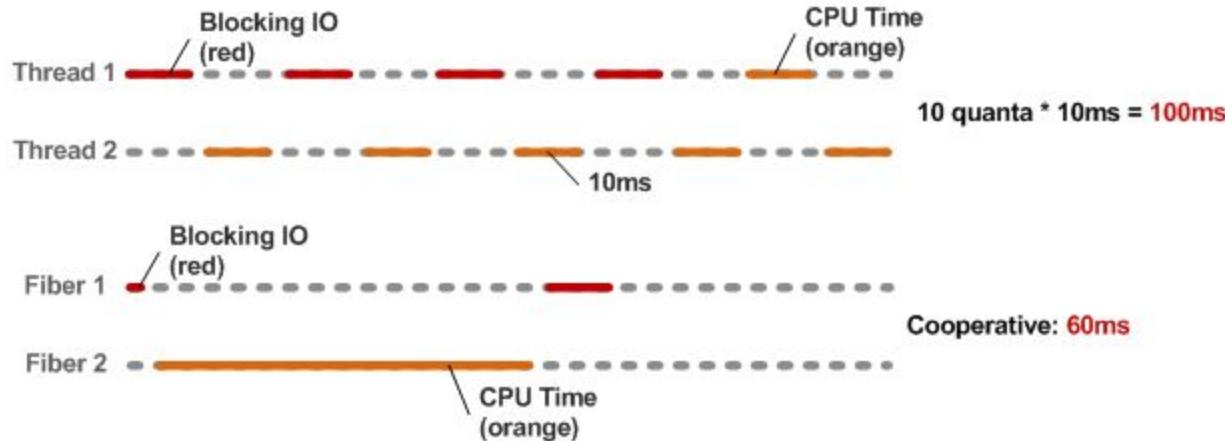
Fiber



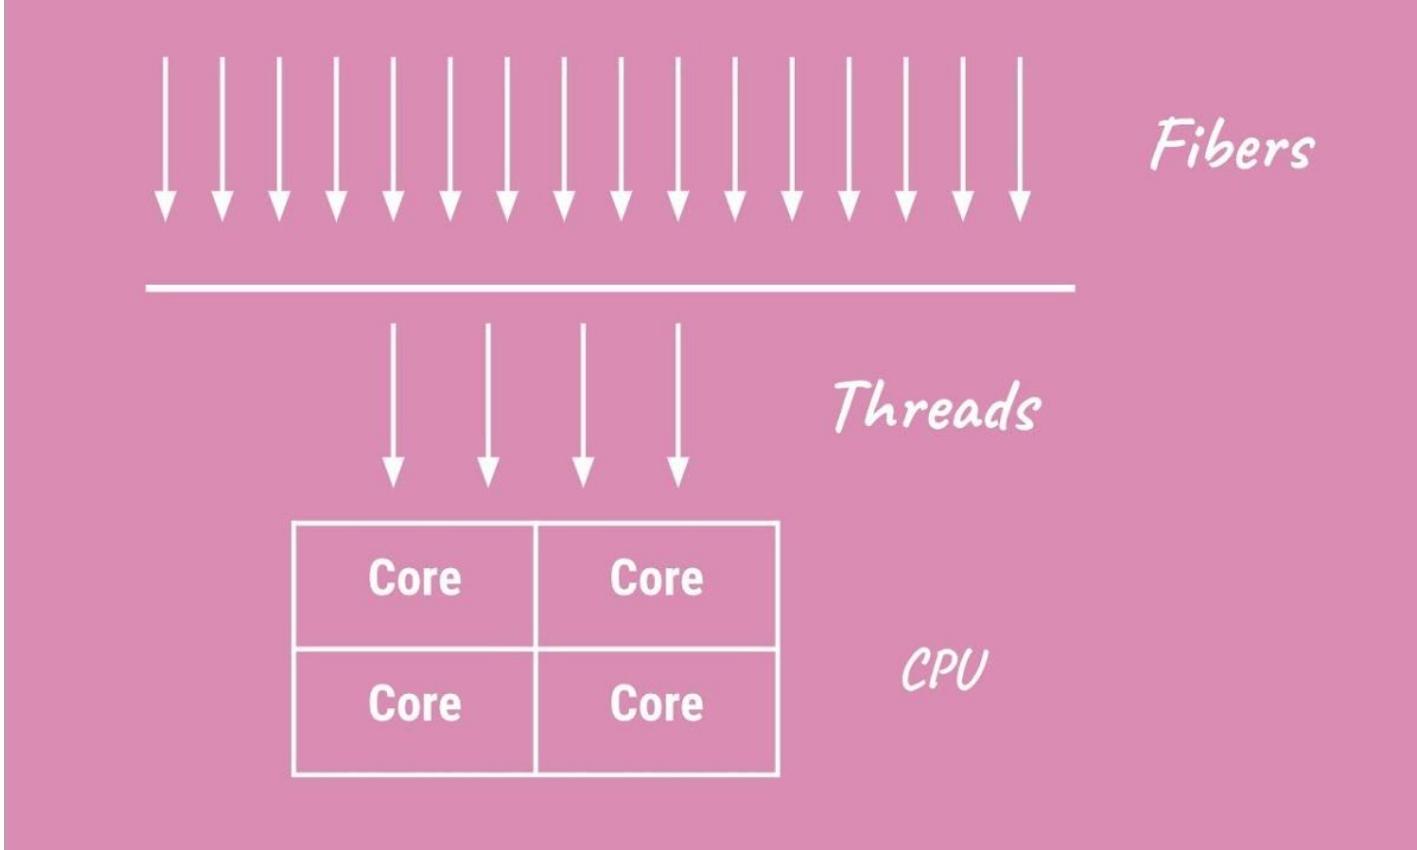
Why do we need fiber

- 1) Efficient use of CPU resources
- 2) Competition Management
- 3) Asynchronous operations
- 4) Parallel algorithms
- 5) Flexibility and control

Fiber scheduling



Why fibers better





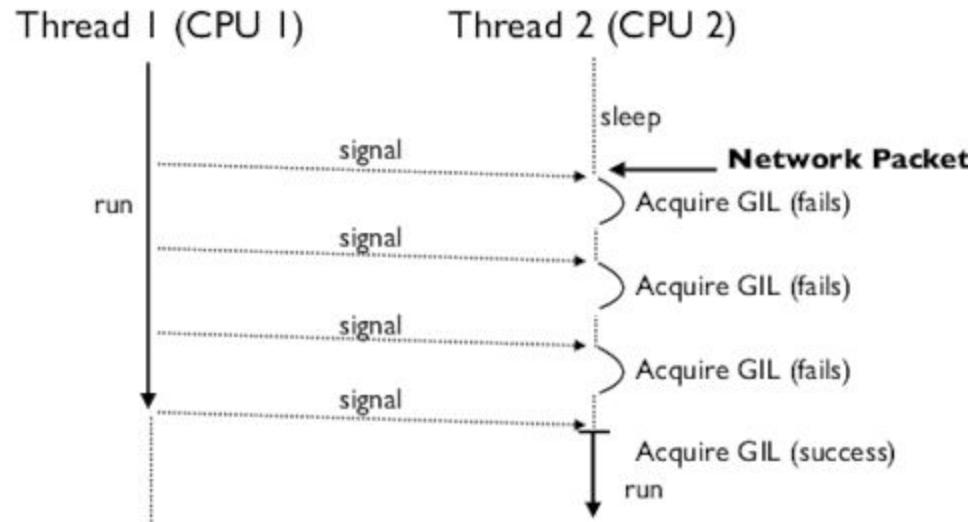
EPIC

Institute of Technology
Powered by epam

07

GIL

GIL

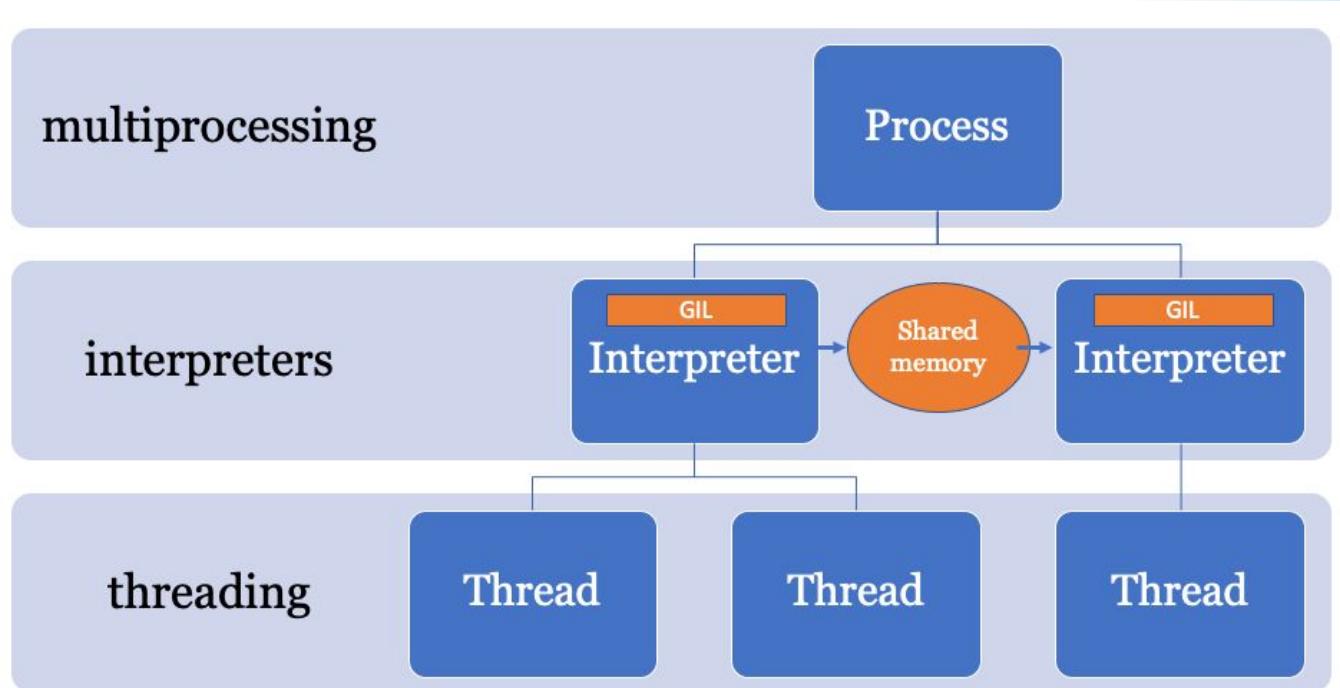




EPIC

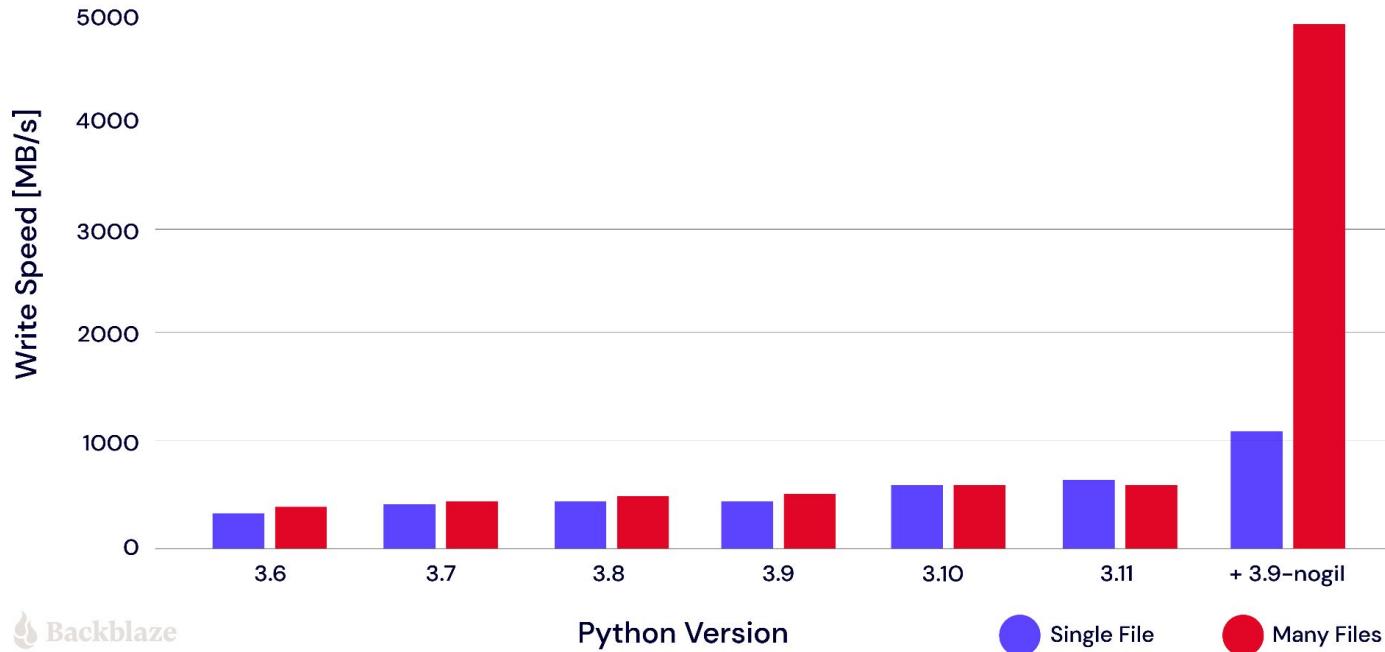
Institute of Technology
Powered by epam

GIL



GIL

B2 CLI Download Speed on Different Python Versions





EPIC

Institute of Technology
Powered by epam

*

Good to know



EPIC

Institute of Technology
Powered by epam

Join

```
#include <iostream>
#include <thread>

// Function to be executed by the thread
void threadFunction() {
    std::cout << "This is a thread." << std::endl;
}

int main() {
    // Creating a thread object and passing it the function to execute
    std::thread myThread(threadFunction);

    std::cout << "This is the main thread." << std::endl;

    // Waiting for the thread to finish
    myThread.join();

    return 0;
}
```



EPIC

Institute of Technology
Powered by epam

Join

A screenshot of a macOS terminal window. The window title bar shows the Apple logo, a house icon, and a tilde icon. The command line shows the path `./a.out`. The main pane displays two lines of text: "This is the main thread." and "This is a thread.".

```
./a.out
This is the main thread.
This is a thread.
```



EPIC

Institute of Technology
Powered by epam

Forget join

```
#include <iostream>
#include <thread>

// Function to be executed by thread
void threadFunction() {
    std::cout << "This is a thread." << std::endl;
}

int main() {
    // Creating a thread object
    std::thread myThread(threadFunction);

    std::cout << "This is the main thread" << std::endl;

    // Forget to wait
    // myThread.join();

    return 0;
}
```



EPIC

Institute of Technology
Powered by epam

Forget join

```
apple ~ ~/Desktop
g++ thread_example.cpp -o a.out

apple ~ ~/Desktop
./a.out

This is the main thread
libc++abi: terminating
This is a thread.
zsh: abort      ./a.out
```



EPIC

Institute of Technology
Powered by epam

Why?

```
~thread() _N0EXCEPT
{ // clean up
    if (joinable())
        XSTD terminate();
}
```



EPIC

Institute of Technology
Powered by epam

Detach

```
#include <iostream>
#include <thread>

// Function to be executed by thread
void threadFunction() {
    std::cout << "This is a thread." << std::endl;
}

int main() {
    // Creating a thread object
    std::thread myThread(threadFunction);

    std::cout << "This is the main thread" << std::endl;

    // Forget to wait
    myThread.detach();

    return 0;
}
```



EPIC

Institute of Technology
Powered by epam

Detach

```
apple ~/Desktop
g++ thread_example.cpp -o a.out

apple ~/Desktop
./a.out
```

This is the main thread



EPIC

Institute of Technology
Powered by epam

Detach

```
#include <iostream>
#include <thread>

// Function to be executed by thread
void threadFunction() {
    std::cout << "This is a thread." << std::endl;
}

int main() {
    // Creating a thread object
    std::thread myThread(threadFunction);

    std::cout << "This is the main thread" << std::endl;

    // Forget to wait
    myThread.detach();
}
```



EPIC

Institute of Technology
Powered by epam

Detach

```
Apple ✚ ~/Desktop
└─ g++ thread_example.cpp -o a.out ┘

Apple ✚ ~/Desktop
└─ ./a.out ┘

This is the main thread
This is a thread.
```

Join

1. Wait for the thread, so work longer
2. Wait for the thread, so work stable

Detach

1. Don't wait for the thread, so work faster
2. Don't wait for the thread, so we need to check that everything is finished

Simple code

```
1. #include <string>
2. #include <thread>
3. #include <iostream>
4. #include <functional>
5.
6. using namespace std;
7.
8. void ChangeCurrentMissileTarget(string& targetCity)
9. {
10.     targetCity = "Metropolis";
11.     cout << " Changing The Target City To " << targetCity << endl;
12. }
13.
14. int main()
15. {
16.     string targetCity = "Star City";
17.     ChangeCurrentMissileTarget(targetCity);
18.
19.     cout << "Current Target City is " << targetCity << endl;
20.
21.     return 0;
22. }
```

Success #stdin #stdout 0.01s 5304KB

 comments (?)

 stdin

 copy

Standard input is empty

 stdout

 copy

Changing The Target City To Metropolis
Current Target City is Metropolis

Multithreading

A screenshot of a code editor interface. The code is a C++ program demonstrating multithreading. It includes includes for string, thread, iostream, and functional. It defines a function ChangeCurrentMissileTarget that changes the target city to "Metropolis". The main function creates a thread t1 to call this function with "Star City" as the argument, then joins with it and prints the current target city.

```
1. #include <string>
2. #include <thread>
3. #include <iostream>
4. #include <functional>
5.
6. using namespace std;
7.
8. void ChangeCurrentMissileTarget(string& targetCity)
9. {
10.     targetCity = "Metropolis";
11.     cout << " Changing The Target City To " << targetCity << endl;
12. }
13.
14. int main()
15. {
16.     string targetCity = "Star City";
17.     thread t1(ChangeCurrentMissileTarget, targetCity);
18.     t1.join();
19.     cout << "Current Target City is " << targetCity << endl;
20.
21.     return 0;
22. }
```

Compilation error #stdin compilation error #stdout 0s 0KB

comments (?)

A screenshot of a terminal or compiler output window. It shows a compilation error for prog.cpp. The error message indicates that the static assertion failed because std::thread arguments must be convertible after conversion to rvalues. The code being compiled is identical to the one shown in the code editor above.

```
compilation info
In file included from prog.cpp:2:
/usr/include/c++/8/thread: In instantiation of 'std::thread::thread(_Callable&&, _Args&& ...) [with _Callable = void (&)(std::__cxx11::basic_string<char>&); _Args = {std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >}]: prog.cpp:17:51:   required from here /usr/include/c++/8/thread:120:17: error: static assertion failed: std::thread arguments must be convertible after conversion to rvalues
```

Correct

```
1. #include <string>
2. #include <thread>
3. #include <iostream>
4. #include <functional>
5.
6. using namespace std;
7.
8. void ChangeCurrentMissileTarget(string& targetCity)
9. {
10.     targetCity = "Metropolis";
11.     cout << " Changing The Target City To " << targetCity << endl;
12. }
13.
14. int main()
15. {
16.     string targetCity = "Star City";
17.     thread t1(ChangeCurrentMissileTarget, std::ref(targetCity));
18.     t1.join();
19.     cout << "Current Target City is " << targetCity << endl;
20.
21.     return 0;
22. }
```

Success #stdin #stdout 0.01s 5304KB

(stdin)

Standard input is empty

(stdout)

Changing The Target City To Metropolis
Current Target City is Metropolis

Possible reasons

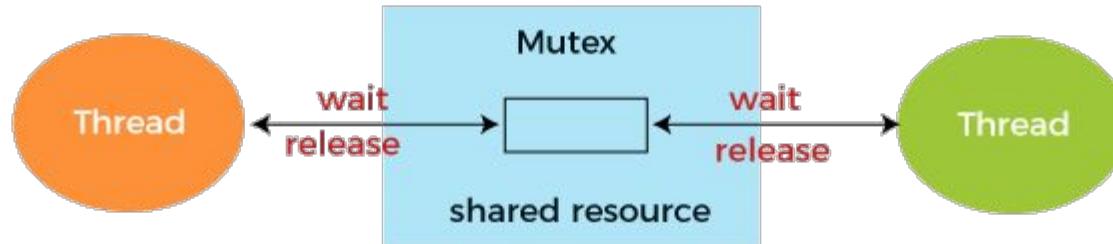
1. Security
2. Easy realization
3. More independence to threads



EPIC

Institute of Technology
Powered by epam

Mutex





EPIC

Institute of Technology
Powered by epam

Mutex usage

```
#include <mutex>
#include <thread>
#include <iostream>

std::mutex mtx; // Create a mutex object

void PrintMessage(const std::string& message)
{
    mtx.lock();
    // Critical section where the thread works with shared data
    std::cout << message << std::endl;
    mtx.unlock();
    // make something without mutex
}

int main()
{
    std::thread t1(PrintMessage, "Hello from Thread 1");
    std::thread t2(PrintMessage, "Hello from Thread 2");

    t1.join();
    t2.join();

    mtx.unlock(); // Release the mutex

    return 0;
}
```



EPIC

Institute of Technology
Powered by epam

Mutex usage

```
└── [apple] ~/Desktop
    g++ thread_example.cpp -o a.out -std=c++17

└── [apple] ~/Desktop
    ./a.out
    Hello from Thread 1
    Hello from Thread 2
```



EPIC

Institute of Technology
Powered by epam

Mistake

```
#include <mutex>
#include <thread>
#include <iostream>

std::mutex mtx; // Create a mutex object

void PrintMessage(const std::string& message)
{
    mtx.lock();
    // Critical section where the thread works with shared data
    std::cout << message << std::endl;
    // mtx.unlock();
    // make something without mutex
}

int main()
{
    std::thread t1(PrintMessage, "Hello from Thread 1");
    std::thread t2(PrintMessage, "Hello from Thread 2");

    t1.join();
    t2.join();

    mtx.unlock(); // Release the mutex

    return 0;
}
```



EPIC

Institute of Technology
Powered by epam

Mistake

```
 MacBook Pro: ~/Desktop
g++ thread_example.cpp -o a.out -std=c++17

MacBook Pro: ~/Desktop
./a.out
Hello from Thread 1
```



EPIC

Institute of Technology
Powered by epam

With guard

```
#include <mutex>
#include <thread>
#include <iostream>

std::mutex mtx; // Create a mutex object

void PrintMessage(const std::string& message)
{
    std::lock_guard<std::mutex> lock(mtx); // Acquire the mutex
    // Critical section where the thread works with shared data
    std::cout << message << std::endl;
    // The mutex is automatically released when the lock_guard goes out of scope
}

int main()
{
    std::thread t1(PrintMessage, "Hello from Thread 1");
    std::thread t2(PrintMessage, "Hello from Thread 2");

    t1.join();
    t2.join();

    return 0;
}
```



EPIC

Institute of Technology
Powered by epam

With guard

```
└── [?] ~/Desktop
    └── g++ thread_example.cpp -o a.out -std=c++17

└── [?] ~/Desktop
    └── ./a.out
        Hello from Thread 2
        Hello from Thread 1
```

Comparison

lock/unlock

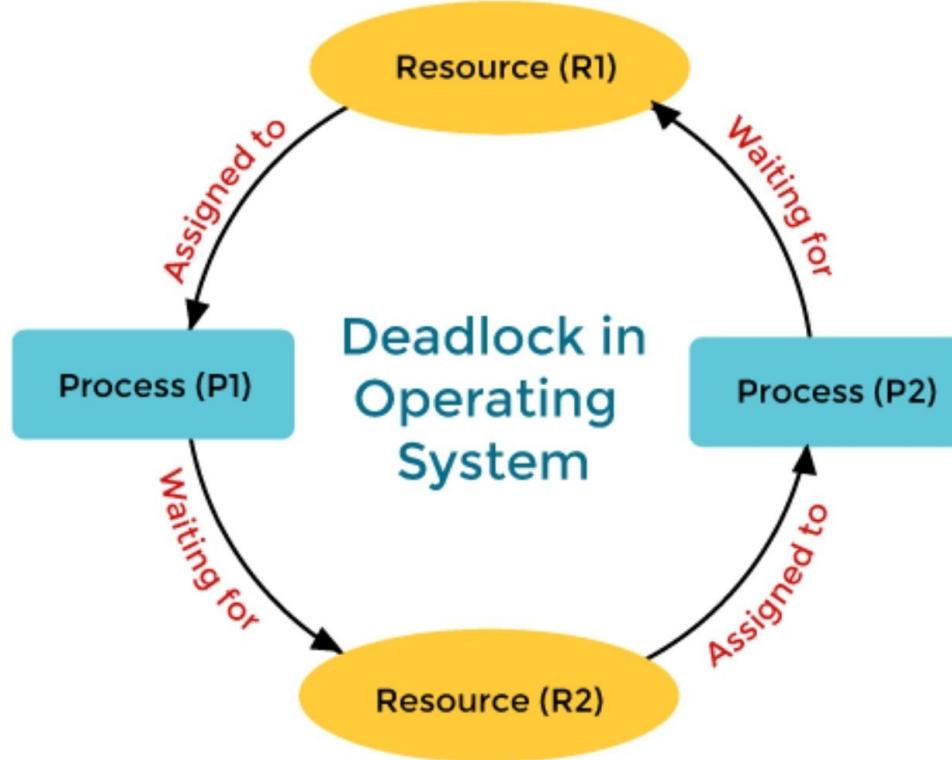
1. Worry about unlock
2. Can work less time because of good unlock

lock_guard

1. Don't worry about unlock
2. Can work more time because of bad unlock



Deadlock





EPIC

Institute of Technology
Powered by epam

Deadlock

```
#include <iostream>
#include <string>
#include <thread>
#include <mutex>

using namespace std;

std::mutex muA;
std::mutex muB;

void CallHome_Th1(string message)
{
    muA.lock();
    std::this_thread::sleep_for(std::chrono::milliseconds(100));
    muB.lock();

    cout << "Thread " << this_thread::get_id() << " says " << message << endl;

    muB.unlock();
    muA.unlock();
}

void CallHome_Th2(string message)
{
    muB.lock();
    std::this_thread::sleep_for(std::chrono::milliseconds(100));
    muA.lock();

    cout << "Thread " << this_thread::get_id() << " says " << message << endl;

    muA.unlock();
    muB.unlock();
}

int main()
{
    thread t1(CallHome_Th1, "Hello from Jupiter");
    thread t2(CallHome_Th2, "Hello from Pluto");

    t1.join();
    t2.join();

    return 0;
}
```



EPIC

Institute of Technology
Powered by epam

Deadlock

```
apple ~ ~/Desktop
g++ thread_example.cpp -o a.out -std=c++17

apple ~ ~/Desktop
./a.out
^C

apple ~ ~/Desktop
```



EPIC

Institute of Technology
Powered by 

That's All Folks!